

# DATA\_WALKTHROUGH

August 27, 2020

## 1 Release One Data Walkthrough

The CRyPTIC data tables and their associated schema is now quite complex. This jupyter notebook walks you through the identity of the tables. It should be read alongside `DATA_SCHEMA.pdf` which you can also find in the `cryptic-tables/` directory. New fields and tables are coloured red.

I'll first explain the tables containing phenotype (i.e. minimum inhibitory concentration) data before considering the genetic data.

This document is available as a PDF or as an interactive jupyter notebook which will let you run each cell containing Python code for yourself. Note that since you have read only access to this folder, if you want to run (i.e. alter) the notebook you will need to copy it to another location on your computer and change `TABLES_PATH` to point to the `cryptic-tables/` folder.

*Philip W Fowler*

27 August 2020

```
[1]: import pandas, numpy
pandas.set_option('display.max_columns', 200)

%matplotlib inline
import matplotlib.pyplot as plt

TABLES_PATH="./"
```

## 2 A note on unique identifiers

CRyPTIC uses a hierarchical set of identifiers.

- `SITEID` a two-digit, left-padded number (`'%02d'`) uniquely identifying each site. Lookup table in `SITES.csv`
- `SUBJID` a string that uniquely identifies each patient in a specific site (might not be unique between sites). Some sites use a left-padded number (e.g. India uses `'%05d'` and China uses `'%04d'`) whilst others are alphanumerical strings. Sometimes country of origin can be inferred.

- **LABID** a string that uniquely identifies the clinical sample. In principle a **SUBJID** can therefore have multiple **LABIDs** associated with, but in practice there is usually only one. Some sites do not need this level of granularity and simply duplicated **SUBJID** in this field.
- **ISOLATENO** integer that identifies the particular isolate tested. Mostly 1. Some cases where the phenotype and genetic data uses different **ISOLATENO**s.
- **SEQ\_REPS** a string associated with the genetics ('sequence repeats'). Again mostly '1', but in principle allows for sequencing to be repeated. Since sequencing appears to mostly fail due to coverage, it is common to see aggregation of short reads e.g. 1\_2\_3
- **READING\_DAY** the number of days of incubation at which the plate was read. Specific to the phenotypic data. **CRyPTIC** uses 14 days unless there is poor growth when it uses 21 days, but there are a wide range of days.

CliRes uses **USUBJID** as its unique identifier which is **SITEID**+"-"+**SUBJID**. Many of the tables use **UNIQUEID** which is a concatenation of the first four identifiers above e.g.

`site.02.subj.0003.lab.20142220007.iso.1`

### 3 Phenotypic data

There are three levels to the phenotype data descending hierarchy in <https://clires2.oucru.org>.

1. **SUBJECT** data. This contains
2. **SAMPLE**
3. **DST** data.

These map onto **SUBJID** i.e. patient-level information, **LABID** i.e. sample-specific information, including other phenotypic tests such as **MGIT** that were run on this sample and lastly **READINGDAY** i.e. the **MICs** read from the 96-well plate after a specified number of days incubation.

All tables were downloaded and populated from <https://clires2.oucru.org> using their WebAPI via the Python **zeep** package. Hence if a **CRyPTIC** lab did not enter their data into CliRes, but instead provided a spreadsheet then they may only has **DST** entries in the **UKMYC\_PHENOTYPES** table and no rows in **SUBJECT** or **SAMPLES**. This applies to CDC Atlanta (01), Italy (06) and Sweden (11) and possibly others.

The original numerical fields have been replaced by descriptive labels to aid interpretation using the **CRyPTIC METADATA FILE SPECIFICATION** . e.g. **GENDER** in **SUBJECTS** contains **MALE/FEMALE/OTHER/UNKNOWN** rather than 1/2/3/9. The version used was April 2018 v2.0. Note that there are some inconsistencies where the spec dictates the values should be 0/1 and they are 1/2.

#### 3.1 SUBJECTS

This only contains 3 fields.

```
[2]: SUBJECTS=pandas.read_pickle(TABLES_PATH+"SUBJECTS.pkl.gz")
SUBJECTS[:3]
```

```
[2]:
```

		GENDER	COUNTRY_OF_ORIGIN
SITEID	SUBJID		
02	0958	male	NaN
	0823	unknown	NaN
	0359	unknown	CHN

```
[3]: SUBJECTS.GENDER.value_counts(dropna=False).sort_index()
```

```
[3]: female      4916
      male       6733
      other         2
      unknown    3848
      Name: GENDER, dtype: int64
```

COUNTRY\_OF\_ORIGIN was not a mandatory field so there are >2,500 missing values.

```
[4]: SUBJECTS.COUNTRY_OF_ORIGIN.value_counts(dropna=False)[:5]
```

```
[4]: IND      5177
      PER      3454
      NaN      3269
      DEU       956
      CHN       896
      Name: COUNTRY_OF_ORIGIN, dtype: int64
```

These are described using the [ISO 3166-1 alpha-3](#) country codes. To help with drawing maps later, there is also a lookup table containing all the 3 letter codes, a proper name and, crucially, the latitude and longitude.

```
[5]: COUNTRIES_LOOKUP=pandas.read_csv("COUNTRIES_LOOKUP.csv")
      COUNTRIES_LOOKUP[:5]
```

```
[5]:
```

	COUNTRY_NAME	COUNTRY_CODE_2_LETTER	COUNTRY_CODE_3_LETTER	\
0	Afghanistan	AF	AFG	
1	Albania	AL	ALB	
2	Algeria	DZ	DZA	
3	American Samoa	AS	ASM	
4	Andorra	AD	AND	

	COUNTRY_CODE_NUMERIC	LAT	LONG
0	4	33.0000	65.0
1	8	41.0000	20.0
2	12	28.0000	3.0
3	16	-14.3333	-170.0
4	20	42.5000	1.6

The COUNTRY\_CODE\_3\_LETTER column (after appropriate renaming) can be used to join to other COUNTRY\_OF\_ORIGIN or COUNTRY\_WHERE\_SAMPLE\_TAKEN to plot maps using (LAT, LONG)

```
[6]: df=SUBJECTS.reset_index()
df=df.
      ↳merge(COUNTRIES_LOOKUP,left_on="COUNTRY_OF_ORIGIN",right_on="COUNTRY_CODE_3_LETTER",how='le
df.COUNTRY_NAME.value_counts()[:5]
```

```
[6]: India      5177
Peru      3454
Germany    956
China      896
Taiwan     392
Name: COUNTRY_NAME, dtype: int64
```

### 3.2 SAMPLES

```
[7]: SAMPLES=pandas.read_pickle(TABLES_PATH+"SAMPLES.pkl.gz")
SAMPLES[:3]
```

```
[7]:
```

			COUNTRY_WHERE_SAMPLE_TAKEN		REGION \
SITEID	SUBJID	LABID			
02	0958	22A197		CHN	CHONGQING
	0823	2013241494		CHN	GUIZHOU
	0359	222018-14		CHN	China ChongQing

			COLLECTION_DATE \
SITEID	SUBJID	LABID	
02	0958	22A197	2017-12-04 00:00:00+07:00
	0823	2013241494	2013-10-06 00:00:00+07:00
	0359	222018-14	2014-01-01 00:00:00+07:00

			ISOLATE_COLLECTED_PROSPECTIVELY	ANATOMICAL_ORIGIN \
SITEID	SUBJID	LABID		
02	0958	22A197	False	not known
	0823	2013241494	False	not known
	0359	222018-14	False	not known

			SMEAR_RESULT	WGS_SEQUENCING_PLATFORM	XPRT_MTB_RIF \
SITEID	SUBJID	LABID			
02	0958	22A197	not known	HiSeq	not tested
	0823	2013241494	not known	HiSeq	not tested
	0359	222018-14	not known	HiSeq	not tested

			HAIN_RIF	HAIN_INH	HAIN_FL	HAIN_AM \
SITEID	SUBJID	LABID				
02	0958	22A197	not tested	not tested	not tested	not tested
	0823	2013241494	not tested	not tested	not tested	not tested
	0359	222018-14	not tested	not tested	not tested	not tested

SITEID	SUBJID	LABID	HAIN_ETH	SMOKER	INJECT_DRUG_USER	IS_HOMELESS	\
02	0958	22A197	not tested	not known	not known	False	
	0823	2013241494	not tested	not known	not known	False	
	0359	222018-14	not tested	not known	not known	False	

SITEID	SUBJID	LABID	IS_IMPRISONED	HIV	DIABETES	WHO_OUTCOME
02	0958	22A197	False	not known	not known	not known
	0823	2013241494	False	not known	not known	not known
	0359	222018-14	False	not known	not known	not known

Usefully COUNTRY\_WHERE\_SAMPLE\_TAKEN is a mandatory field. REGION is freeform text so will require cleaning if it is to be used.

```
[8]: SAMPLES.COUNTRY_WHERE_SAMPLE_TAKEN.value_counts(dropna=False)[:6]
```

```
[8]: IND      5109
PER      3450
ZAF      2267
CHN      1509
VNM      1112
DEU       851
Name: COUNTRY_WHERE_SAMPLE_TAKEN, dtype: int64
```

The COLLECTION\_DATE might be useful, although beware samples from 1900 and 2201!

```
[9]: SAMPLES[['COLLECTION_DATE']].groupby(SAMPLES.COLLECTION_DATE.dt.year).count()
```

```
[9]: COLLECTION_DATE
COLLECTION_DATE
1900              2
1986              1
2001              1
2003             30
2004             31
2005             13
2006              3
2007            111
2008             41
2009            133
2010             83
2011            122
2012            495
2013           1360
2014            856
2015            923
```

2016	1011
2017	2825
2018	4540
2019	2921
2020	5
2201	1

```
[10]: SAMPLES.ISOLATE_COLLECTED_PROSPECTIVELY.value_counts(dropna=False)
```

```
[10]: False    12451
      True     3057
      Name: ISOLATE_COLLECTED_PROSPECTIVELY, dtype: int64
```

As one might expect the vast majority of samples are respiratory. Again this was not a required field so beware the large number of 'not known' values.

```
[11]: SAMPLES.ANATOMICAL_ORIGIN.value_counts(dropna=False).
      ↪sort_values(ascending=False)
```

```
[11]: Respiratory    7393
      not known    5306
      Other known site 1843
      Lymph node    461
      CSF           344
      Pleural       134
      Non-respiratory, site not known 22
      Bone          5
      Name: ANATOMICAL_ORIGIN, dtype: int64
```

There is also some smear data for some samples

```
[12]: SAMPLES.SMEAR_RESULT.value_counts(dropna=False)
```

```
[12]: not known    6388
      Negative    3351
      +           2619
      +++         1340
      ++          1223
      Scanty       587
      Name: SMEAR_RESULT, dtype: int64
```

You can also check what sequencing platform was used (although this is, perhaps more correctly, recorded in the metadata spreadsheets sent to the EBI along with the FASTQ files)

```
[13]: SAMPLES.WGS_SEQUENCING_PLATFORM.value_counts(dropna=False)
```

```
[13]: NextSeq    6706
      HiSeq     5499
```

```

Other          2173
MiSeq          969
NovaSeq6000    161
Name: WGS_SEQUENCING_PLATFORM, dtype: int64

```

Some samples have been on an Xpert MTB/RIF cartridge (note that there were only a handful of Ultra samples so these were discarded)

```
[14]: SAMPLES.XPERT_MTB_RIF.value_counts(dropna=False)
```

```

[14]: not tested          13859
      RIF susceptible      683
      test inconclusive    482
      RIF resistant        480
      test failed          4
      Name: XPERT_MTB_RIF, dtype: int64

```

A similar number also have Hain LPA results recorded for RIF, INH,

```
[15]: SAMPLES.HAIN_RIF.value_counts()
```

```

[15]: not tested          13983
      susceptible          930
      resistant           453
      test inconclusive    140
      test failed          2
      Name: HAIN_RIF, dtype: int64

```

```
[16]: SAMPLES.HAIN_INH.value_counts()
```

```

[16]: not tested          13983
      susceptible          865
      resistant           518
      test inconclusive    140
      test failed          2
      Name: HAIN_INH, dtype: int64

```

```
[17]: SAMPLES.HAIN_ETH.value_counts()
```

```

[17]: not tested          15331
      test inconclusive    140
      susceptible          30
      resistant            7
      Name: HAIN_ETH, dtype: int64

```

The below is understood to be a generic fluoroquinolone result

```
[18]: SAMPLES.HAIN_FL.value_counts()
```

```
[18]: not tested      15331
      test inconclusive  140
      susceptible      22
      resistant        15
      Name: HAIN_FL, dtype: int64
```

The below is understood to be a generic aminoglycoside result

```
[19]: SAMPLES.HAIN_AM.value_counts()
```

```
[19]: not tested      15331
      test inconclusive  140
      susceptible      32
      resistant         5
      Name: HAIN_AM, dtype: int64
```

Now we have some sparse lifestyle data. First is whether they smoked or not.

```
[20]: SAMPLES.SMOKER.value_counts()
```

```
[20]: not known      14260
      no              937
      yes, currently  172
      yes, previously 139
      Name: SMOKER, dtype: int64
```

```
[21]: SAMPLES.INJECT_DRUG_USER.value_counts()
```

```
[21]: not known      14268
      no              1104
      yes, previously  113
      yes, currently   23
      Name: INJECT_DRUG_USER, dtype: int64
```

```
[22]: SAMPLES.IS_HOMELESS.value_counts()
```

```
[22]: False      15490
      True        18
      Name: IS_HOMELESS, dtype: int64
```

```
[23]: SAMPLES.IS_IMPRISONED.value_counts()
```

```
[23]: False      15438
      True        70
      Name: IS_IMPRISONED, dtype: int64
```

```
[24]: SAMPLES.HIV.value_counts()
```



```
[24]: not known          12895
      tested, negative    1938
      tested, positive    669
      not tested          6
      Name: HIV, dtype: int64
```

```
[25]: SAMPLES.DIABETES.value_counts()
```

```
[25]: not known          13421
      tested, not diabetic  1892
      tested, type 2 diabetes  107
      tested, unknown type    68
      tested, type 1 diabetes  20
      Name: DIABETES, dtype: int64
```

Finally, a small number had a WHO outcome field recorded. These have been translated

```
[26]: SAMPLES.WHO_OUTCOME.value_counts(dropna=False).sort_index()
```

```
[26]: cured              1567
      died                277
      lost to follow-up or defaulted  265
      not evaluated        29
      not known           12531
      treatment completed    739
      treatment failed       100
      Name: WHO_OUTCOME, dtype: int64
```

### 3.3 UKMYC\_PLATES

UKMYC\_PLATES contains one row per plate. It is a simplified view of the old PLATES and PLATE\_MEASUREMENTS tables and hence contains ‘the’ reading and therefore the READINGDAY, which in most cases will be day 14. All other readings taken on other reading days are not shown in this view.

```
[27]: UKMYC_PLATES=pandas.read_pickle(TABLES_PATH+"UKMYC_PLATES.pkl.gz")
      UKMYC_PLATES[:3]
```

```
[27]:
```

	SITEID	SUBJID	LABID	ISOLATENO	\
UNIQUEID					
site.11.subj.MDR044.lab.SWE-33.iso.1	11	MDR044	SWE-33	1	
site.11.subj.MDR045.lab.SWE-34.iso.1	11	MDR045	SWE-34	1	
site.11.subj.MDR046.lab.SWE-35.iso.1	11	MDR046	SWE-35	1	

	READINGDAY	BELONGS_GPI	PLATEDESIGN	\
UNIQUEID				
site.11.subj.MDR044.lab.SWE-33.iso.1	10	False	UKMYC5	

site.11.subj.MDR045.lab.SWE-34.iso.1	10	True	UKMYC5
site.11.subj.MDR046.lab.SWE-35.iso.1	10	True	UKMYC5

TREE\_PATH \

UNIQUEID

site.11.subj.MDR044.lab.SWE-33.iso.1	dat/CRyPTIC2/V2/11/MDR044/SWE-33/1/10/
site.11.subj.MDR045.lab.SWE-34.iso.1	dat/CRyPTIC2/V2/11/MDR045/SWE-34/1/10/
site.11.subj.MDR046.lab.SWE-35.iso.1	dat/CRyPTIC2/V2/11/MDR046/SWE-35/1/10/

IMAGEFILENAME IMAGE\_MD5SUM \

UNIQUEID

site.11.subj.MDR044.lab.SWE-33.iso.1	NaN	NaN
site.11.subj.MDR045.lab.SWE-34.iso.1	NaN	NaN
site.11.subj.MDR046.lab.SWE-35.iso.1	NaN	NaN

DUPLICATED\_IMAGE IM\_IMAGE\_DOWNLOADED \

UNIQUEID

site.11.subj.MDR044.lab.SWE-33.iso.1	False	False
site.11.subj.MDR045.lab.SWE-34.iso.1	False	False
site.11.subj.MDR046.lab.SWE-35.iso.1	False	False

IM\_IMAGE\_FILTERED IM\_WELLS\_IDENTIFIED \

UNIQUEID

site.11.subj.MDR044.lab.SWE-33.iso.1	False	False
site.11.subj.MDR045.lab.SWE-34.iso.1	False	False
site.11.subj.MDR046.lab.SWE-35.iso.1	False	False

IM\_POS1GROWTH IM\_POS2GROWTH \

UNIQUEID

site.11.subj.MDR044.lab.SWE-33.iso.1	0.0	0.0
site.11.subj.MDR045.lab.SWE-34.iso.1	0.0	0.0
site.11.subj.MDR046.lab.SWE-35.iso.1	0.0	0.0

IM\_POS\_AVERAGE \

UNIQUEID

site.11.subj.MDR044.lab.SWE-33.iso.1	NaN
site.11.subj.MDR045.lab.SWE-34.iso.1	NaN
site.11.subj.MDR046.lab.SWE-35.iso.1	NaN

IM\_DRUGS\_INCONSISTENT\_GROWTH \

UNIQUEID

site.11.subj.MDR044.lab.SWE-33.iso.1	NaN
site.11.subj.MDR045.lab.SWE-34.iso.1	NaN
site.11.subj.MDR046.lab.SWE-35.iso.1	NaN

TRUST\_PHENOTYPES

UNIQUEID

```

site.11.subj.MDR044.lab.SWE-33.iso.1      True
site.11.subj.MDR045.lab.SWE-34.iso.1      True
site.11.subj.MDR046.lab.SWE-35.iso.1      True

```

```
[28]: UKMYC_PLATES.READINGDAY.value_counts(dropna=False)
```

```

[28]: 14      19834
      21        723
      10         80
      28          0
      7          0
      Name: READINGDAY, dtype: int64

```

The BELONGS\_GPI field is new, and tells us if this plate belong to the ‘Geno-Pheno-Intersection’ i.e. whether it was sequenced and successfully processed using Clockwork by the EBI.

```
[29]: UKMYC_PLATES.BELONGS_GPI.value_counts()
```

```

[29]: True      15039
      False    5598
      Name: BELONGS_GPI, dtype: int64

```

This is a bit less than the original number in the GPI of 15,211 since invalid plates (poor growth, contamination, problems with the control wells) have been excluded.

The PLATEDESIGN field is important since it tells us which antibiotics where on the plate, where they are located and their concentrations. You can look this up via

```
[30]: PLATE_DESIGN=pandas.read_csv(TABLES_PATH+"PLATE_LAYOUT.csv.gz")
      PLATE_DESIGN[:3]
```

```

[30]:  PLATEDESIGN  DRUG  DILUTION  CONC  ROW  COL  BINARY_PHENOTYPE
0      UKMYC5  AMI          7  >8  NaN  NaN              R
1      UKMYC5  AMI          6   8  1.0  1.0              R
2      UKMYC5  AMI          5   4  2.0  1.0              R

```

Note that, for ease, the binary phenotype for each MIC according to the ‘current’ CRyPTIC ECOFFs is also included in this table. The CRyPTIC ECOFFs may change slightly will change the BINARY\_PHENOTYPE assignments – you’ll be notified if this happens. We shall use this later

A 3-letter code is used in all the tables to identify all the drugs. Whilst these are mostly standard/obvious, there is a further lookup table you can use to get more information on what drug each 3-letter code describes.

```
[31]: DRUG_LOOKUP=pandas.read_csv(TABLES_PATH+'DRUG_LOOKUP.csv.gz')
      DRUG_LOOKUP[:3]
```

```

[31]:  SAMPLES_COLUMN  OTHER_PHENOTYPES_COLUMN  DRUG_ABBREVIATION
0      RIFAMPICIN                        RIF              RIF

```

1	RIFMETHOD	RIF_METHOD	RIF
2	RIFOTH	RIF_OTHER	RIF

Note that CRyPTIC Release One is about 60% UKMYC6, at least when it comes to “samples on plates”

```
[32]: UKMYC_PLATES.PLATEDESIGN.value_counts(dropna=False)
```

```
[32]: UKMYC6      12672
      UKMYC5       7965
      Name: PLATEDESIGN, dtype: int64
```

Different to before, this table contains additional fields that allow you to retrieve raw files **directly** from the sharded data tree (assuming you have access). Central to this is the `TREE_PATH` which gives you the relative path to the leaf where the files (in this case images) are stored. For example, if you wanted to retrieve a list of raw images you could do (bit clunky but works)

```
[33]: def return_good_images(row):
      print('/well/bag/pfowler/cryptic/
      ↳'+str(row['TREE_PATH'])+str(row['IMAGEFILENAME'])+'-'+row['PLATEDESIGN']+'-growth.
      ↳jpg')

      GOOD_PLATE_IMAGES=UKMYC_PLATES.loc[(UKMYC_PLATES.IMAGEFILENAME.notna()) &
      ↳(UKMYC_PLATES.IM_IMAGE_FILTERED.notna()) & (~UKMYC_PLATES.DUPLICATED_IMAGE)
      ↳& (UKMYC_PLATES.TRUST_PHENOTYPES)]

      a=GOOD_PLATE_IMAGES[:3].apply(return_good_images,axis=1)
```

```
/well/bag/pfowler/cryptic/dat/CRyPTIC2/V2/01/DR0013/DR0013/1/14/01-DR0013-DR0013
-1-14-UKMYC6-growth.jpg
/well/bag/pfowler/cryptic/dat/CRyPTIC2/V2/01/DR0018/DR0018/1/14/01-DR0018-DR0018
-1-14-UKMYC6-growth.jpg
/well/bag/pfowler/cryptic/dat/CRyPTIC2/V2/01/DR0025/DR0025/1/14/01-DR0025-DR0025
-1-14-UKMYC6-growth.jpg
```

```
[34]: pandas.crosstab(UKMYC_PLATES.PLATEDESIGN,UKMYC_PLATES.BELONGS_GPI)
```

```
[34]: BELONGS_GPI  False  True
      PLATEDESIGN
      UKMYC5       732   7233
      UKMYC6      4866   7806
```

The ratio of plate designs in the GPI is more like 50:50 since we haven’t received the FASTQ files for some of the newer UKMYC6 samples.

The MD5SUM of the image is recorded here so duplicates can be identified. (Note that if `True` all of the image related measurements are discarded and therefore these measurements can never have a `PHENOTYPE_QUALITY` of `HIGH`).

```
[35]: UKMYC_PLATES.DUPLICATED_IMAGE.value_counts()
```

```
[35]: False    20174  
      True      463  
      Name: DUPLICATED_IMAGE, dtype: int64
```

Sometimes the images are duplicated across READINGDAYS for the same sample, which is lazy but not too bad. However, there are also many instances where the same image has been associated with different samples.

```
[36]: UKMYC_PLATES[UKMYC_PLATES.  
      ↪IMAGE_MD5SUM=="7e0d4fce8ecc9f2c9c08f87098c3c85f"] [["DUPLICATED_IMAGE"]]
```

```
[36]:
```

	DUPLICATED_IMAGE
UNIQUEID	
site.04.subj.00033.lab.628880.iso.1	True
site.04.subj.00200.lab.634474.iso.1	True
site.04.subj.01246.lab.719263.iso.1	True


The IM\_IMAGE\_DOWNLOADED, IM\_IMAGE\_FILTERED and IM\_WELLS\_IDENTIFIED boolean fields tell you, respectively, if an image was downloaded, was successfully filtered by AMyGDA and whether 96 (and only 96!) wells were identified by AMyGDA. The last can fail if the image is improperly cropped or if the photo quality is so poor (e.g. washed out) that the algorithm cannot find the edges of the wells.

Note that this means having an image present, as indicated by IMAGEFILENAME does not guarantee that AMyGDA was also able to read it.

If AMyGDA was able to read the plate, then the growth in the two control wells (and their average for convenience) is recorded in the next 3 fields. As shown below the growth in the two control wells is correlated, but also truncated.


```
[37]: UKMYC_PLATES.plot.  
      ↪scatter(x='IM_POS1GROWTH',y='IM_POS2GROWTH',figsize=(8,8),marker='.')
```

```
[37]: <matplotlib.axes._subplots.AxesSubplot at 0x12c99f190>
```



```
[38]: UKMYC_PLATES.IM_POS_AVERAGE.plot.hist(figsize=(8,4),bins=50)
```


```
[38]: <matplotlib.axes._subplots.AxesSubplot at 0x12d5a9290>
```



The IM\_DRUGS\_INCONSISTENT\_GROWTH field is the total number of drugs on this plate that AMyGDA was unable to read due to skipped wells etc. In theory a high number is a hint that this is a difficult plate to read. By definition, if AMyGDA cannot read the image, then that measurement will have been sent to BashTheBug for reading.

```
[39]: UKMYC_PLATES.IM_DRUGS_INCONSISTENT_GROWTH.plot.  
      ↪hist(figsize=(8,4),bins=range(0,14))
```

```
[39]: <matplotlib.axes._subplots.AxesSubplot at 0x12d07c410>
```



Finally, the TRUST\_PHENOTYPES column indicates which samples it is suspected are subject to

systematic measurement error and therefore need to be excluded from the UKMYC\_PHENOTYPES table. At present these samples are excluded \* all of SITEID=='13' (very high number of discrepant recorded by Taiwan) \* 02350>=SUBJID>='01575' for SITEID=='04' (poor growth episode flagged by India) Note that SUBJID>02350 for SITEID==04 contains an elevated number of discrepancies and you may wish to also exclude these 1797 samples and see what the effect is on your analysis.

```
[40]: UKMYC_PLATES.loc[~UKMYC_PLATES.TRUST_PHENOTYPES][:3]
```

```
[40]:
```

	SITEID	SUBJID	LABID	ISOLATENO	\
UNIQUEID					
site.04.subj.01575.lab.728967.iso.1	04	01575	728967	1	
site.04.subj.01576.lab.729539.iso.1	04	01576	729539	1	
site.04.subj.01577.lab.725940.iso.1	04	01577	725940	1	

	READINGDAY	BELONGS_GPI	PLATEDESIGN	\
UNIQUEID				
site.04.subj.01575.lab.728967.iso.1	14	True	UKMYC5	
site.04.subj.01576.lab.729539.iso.1	14	True	UKMYC5	
site.04.subj.01577.lab.725940.iso.1	14	True	UKMYC5	

	TREE_PATH	\
UNIQUEID		
site.04.subj.01575.lab.728967.iso.1	dat/CRyPTIC2/V2/04/01575/728967/1/14/	
site.04.subj.01576.lab.729539.iso.1	dat/CRyPTIC2/V2/04/01576/729539/1/14/	
site.04.subj.01577.lab.725940.iso.1	dat/CRyPTIC2/V2/04/01577/725940/1/14/	

	IMAGEFILENAME	\
UNIQUEID		
site.04.subj.01575.lab.728967.iso.1	04-01575-728967-1-14	
site.04.subj.01576.lab.729539.iso.1	04-01576-729539-1-14	
site.04.subj.01577.lab.725940.iso.1	04-01577-725940-1-14	

	IMAGE_MD5SUM	\
UNIQUEID		
site.04.subj.01575.lab.728967.iso.1	98e488b15b40f09bff1b84dbd3327d76	
site.04.subj.01576.lab.729539.iso.1	57cd0d552bd53b48d7152428ada07c77	
site.04.subj.01577.lab.725940.iso.1	4ce81fb9bf44d56d3b1fef2e0e504293	

	DUPLICATED_IMAGE	IM_IMAGE_DOWNLOADED	\
UNIQUEID			
site.04.subj.01575.lab.728967.iso.1	False	True	
site.04.subj.01576.lab.729539.iso.1	False	True	
site.04.subj.01577.lab.725940.iso.1	False	True	

	IM_IMAGE_FILTERED	IM_WELLS_IDENTIFIED	\
UNIQUEID			
site.04.subj.01575.lab.728967.iso.1	True	True	



site.04.subj.01576.lab.729539.iso.1	True	True
site.04.subj.01577.lab.725940.iso.1	True	True

	IM_POS1GROWTH	IM_POS2GROWTH	\
UNIQUEID			
site.04.subj.01575.lab.728967.iso.1	28.11	26.92	
site.04.subj.01576.lab.729539.iso.1	7.62	10.57	
site.04.subj.01577.lab.725940.iso.1	9.48	23.15	

	IM_POS_AVERAGE	\
UNIQUEID		
site.04.subj.01575.lab.728967.iso.1	27.52	
site.04.subj.01576.lab.729539.iso.1	9.10	
site.04.subj.01577.lab.725940.iso.1	16.31	

	IM_DRUGS_INCONSISTENT_GROWTH	\
UNIQUEID		
site.04.subj.01575.lab.728967.iso.1		0.0
site.04.subj.01576.lab.729539.iso.1		0.0
site.04.subj.01577.lab.725940.iso.1		2.0

	TRUST_PHENOTYPES
UNIQUEID	
site.04.subj.01575.lab.728967.iso.1	False
site.04.subj.01576.lab.729539.iso.1	False
site.04.subj.01577.lab.725940.iso.1	False

Although we start off with 15,211 in the GPI, we lose some since the plates are not readable. If we also exclude those which are under investigation due to high levels of discrepancies, then we reach 14,159!

```
[41]: pandas.crosstab(UKMYC_PLATES.TRUST_PHENOTYPES, UKMYC_PLATES.BELONGS_GPI)
```

```
[41]: BELONGS_GPI      False   True
      TRUST_PHENOTYPES
      False           62    880
      True           5536  14159
```

Because most of the plates excluded by TRUST\_PHENOTYPES are UKMYC5, we end up at 45:55% for UKMYC5/6.

```
[42]: UKMYC_PLATES.loc[(UKMYC_PLATES.TRUST_PHENOTYPES) & (UKMYC_PLATES.BELONGS_GPI)] .
      ↪ PLATEDESIGN.value_counts()
```

```
[42]: UKMYC6      7804
      UKMYC5      6355
      Name: PLATEDESIGN, dtype: int64
```

Of these, 12,984 (92%) have images.

```
[43]: len(UKMYC_PLATES.loc[(UKMYC_PLATES.TRUST_PHENOTYPES) & (UKMYC_PLATES.  
    ↳ BELONGS_GPI) & (UKMYC_PLATES.IM_POS_AVERAGE.notna())])
```

```
[43]: 12984
```

### 3.4 UKMYC\_PHENOTYPES

This is the **core** phenotype table. It aggregates and summarises all the above and presents the current ‘best valid reading’ on each DRUG for each UNIQUEID.

```
[44]: UKMYC_PHENOTYPES=pandas.read_pickle(TABLES_PATH+"UKMYC_PHENOTYPES.pkl.gz")  
UKMYC_PHENOTYPES[:3]
```

```
[44]:
```

UNIQUEID	DRUG	PLATEDESIGN	BELONGS_GPI	SITEID	\
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	UKMYC5	False	06	
site.04.subj.00861.lab.713588.iso.1	ETH	UKMYC5	True	04	
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	UKMYC5	True	08	

UNIQUEID	DRUG	DILUTION	PHENOTYPE_QUALITY	\
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	5.0	MEDIUM	
site.04.subj.00861.lab.713588.iso.1	ETH	2.0	HIGH	
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	1.0	HIGH	

UNIQUEID	DRUG	READINGDAY	PRIMARY_DILUTION	\
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	14	5.0	
site.04.subj.00861.lab.713588.iso.1	ETH	14	2.0	
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	14	1.0	

UNIQUEID	DRUG	PRIMARY_METHOD	AMYGDA_DILUTION	\
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	VZ	NaN	
site.04.subj.00861.lab.713588.iso.1	ETH	VZ	2.0	
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	VZ	1.0	

UNIQUEID	DRUG	BASHTHEBUG_DILUTION	\
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	NaN	
site.04.subj.00861.lab.713588.iso.1	ETH	3.0	
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	1.0	


UNIQUEID	DRUG	BASHTHEBUGPRO_DILUTION	\
----------	------	------------------------	---

site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	NaN
site.04.subj.00861.lab.713588.iso.1	ETH	NaN
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	NaN

UNIQUEID	DRUG	PHENOTYPE_DESCRIPTION \
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	VZ ONLY
site.04.subj.00861.lab.713588.iso.1	ETH	VZ,IM AGREE
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	VZ,IM AGREE

UNIQUEID	DRUG	BASHTHEBUG_NUMBER_CLASSIFICATIONS \
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	NaN
site.04.subj.00861.lab.713588.iso.1	ETH	11.0
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	11.0

UNIQUEID	DRUG	MIC	LOG2MIC	BINARY_PHENOTYPE
site.06.subj.CL4441.lab.06MIL0824.iso.1	LZD	0.5	-1.00	S
site.04.subj.00861.lab.713588.iso.1	ETH	0.5	-1.00	S
site.08.subj.24TB00-059.lab.2444.iso.1	RFB	<=0.06	-4.06	S



Each reading flows down the above quality assurance process; first a reading (PRIMARY\_DILUTION) is recorded in the laboratory by the scientist using the PRIMARY\_METHOD. This is almost always VZ i.e. Vizion, but some labs were only able to use Mirrored Box (MB) for some measurements. This is almost always after 14 days of incubation (READINGDAY==14), but if a reading was not possible, we then consider the day 21 reading, if available.

Note that we use negative DILUTIONs to indicate there was a problem with that reading (the MIC will be NaN).

- -1 cannot read this particular drug for some reason (but usually can read the others off a plate). For AMyGDA (IM) this indicates the presence of one or more skip wells.
- -2 no or insufficient growth in one or both of the control wells

If no photograph of the plate was stored, or the image was subsequently duplicated (as indicated in UKMYC\_PLATES), then only one reading is possible, the PRIMARY\_DILUTION as done by the PRIMARY\_METHOD. In these cases, the PHENOTYPE\_QUALITY is left as the default, which is MEDIUM.

About 10.3% of all readings have no image.

```
[45]: len(UKMYC_PHENOTYPES.loc[UKMYC_PHENOTYPES.PHENOTYPE_QUALITY=="MEDIUM"])/  
      ↳ len(UKMYC_PHENOTYPES)
```

```
[45]: 0.1027750948734134
```

The remaining measurements do have a (not duplicated) image. Each image is analysed by AMyGDA and the dilution recorded in AMYGDA\_DILUTION. In about 54% of cases, this exactly agrees with the PRIMARY\_DILUTION and hence this measurement is marked as PHENOTYPE\_QUALITY='HIGH'.

The remaining 46% are sent to BashTheBug for assessment by citizen scientists. Once BASHTHEBUG\_NUMBER\_CLASSIFICATIONS>=11 the media value is returned as the consensus and populated in BASHTHEBUG\_DILUTION.

Of these, in about

- 38% of cases, BASHTHEBUG\_DILUTION and PRIMARY\_DILUTION are identical, suggesting that AMyGDA incorrectly read the plate (due to e.g. low growth or artefacts).
- 12% of cases BASHTHEBUG\_DILUTION and AMYGDA\_DILUTION are identical, suggesting that the laboratory scientist made a measurement or data entry error.
- 50% of cases, all three measurements are different.

If two measurements exactly agree, then the measurement is marked as PHENOTYPE\_QUALITY='HIGH', otherwise if all three disagree, then PHENOTYPE\_DESCRIPTION='ALL DISAGREE' and the PHENOTYPE\_QUALITY is marked as LOW.

It is recommended that, unless you have a good reason to the contrary, to only use readings where PHENOTYPE\_QUALITY is HIGH. We would be very interested in knowing what, if any, the effect of this QA workflow is, so would also be interested in seeing the effect of ignoring the PHENOTYPE\_QUALITY i.e. just using the PRIMARY\_DILUTION.

```
[46]: UKMYC_PHENOTYPES.PHENOTYPE_DESCRIPTION.value_counts()
```

```
[46]: VZ,IM AGREE      128038  
      ALL DISAGREE   54419  
      VZ,BB AGREE    40726  
      VZ ONLY        26684  
      BB,IM AGREE    12771  
      BB RUNNING      344  
      Name: PHENOTYPE_DESCRIPTION, dtype: int64
```

```
[47]: pandas.crosstab(UKMYC_PHENOTYPES.PHENOTYPE_DESCRIPTION,UKMYC_PHENOTYPES.  
      ↳ PHENOTYPE_QUALITY,margins=True)
```

```
[47]: PHENOTYPE_QUALITY      HIGH      LOW  MEDIUM      All
      PHENOTYPE_DESCRIPTION
      ALL DISAGREE           0  54419           0  54419
      BB RUNNING             0       0       344    344
      BB,IM AGREE          12771       0       0  12771
      VZ ONLY               0       0  26684  26684
      VZ,BB AGREE          40726       0       0  40726
      VZ,IM AGREE          128038       0       0  128038
      All                   181535  54419  27028  262982
```

As mentioned above, the 344 rows where BashTheBug does not appear to have finished are glitches and will remain for the time being at least. 293 are for PAS which is excluded in all analyses.

That leaves 26,684 (10% of total) measurements where there is no image (or the image was a duplicate) and so only one measurement (usually VZ) is possible and hence these cannot progress any further than a PHENOTYPE\_QUALITY of MEDIUM.

Of the remaining 233,948 measurements, 77% have two or more measurement methods (VZ/IM/BB) in exact concordance and therefore are classified as having a HIGH PHENOTYPE\_QUALITY. In 70.8% of these, the AMyGDA measurement agreed with Vizion whilst in 22.2% and 7.0% of cases BashTheBug agreed with Vizion or AMyGDA, respectively. The latter set contain mistakes made by the laboratory scientist and therefore is an upper estimate of the laboratory reading error rate.

```
[48]: UKMYC_PHENOTYPES.loc[UKMYC_PHENOTYPES.PHENOTYPE_QUALITY=='HIGH'] .
      ↳PHENOTYPE_DESCRIPTION.value_counts(normalize=True)
```

```
[48]: VZ,IM AGREE      0.705308
      VZ,BB AGREE     0.224342
      BB,IM AGREE     0.070350
      VZ ONLY         0.000000
      BB RUNNING      0.000000
      ALL DISAGREE    0.000000
      Name: PHENOTYPE_DESCRIPTION, dtype: float64
```

There are 53,129 rows where all three methods disagree. The last time the analysis was run 79% of these have been processed by BashTheBugPRO. (The figure stands now at 90% - 12 Aug 2020)

In future, the number of HIGH quality measurements will be increased by allowing their consensus to overrule i.e. they will arbitrate. They may choose to overrule the three methods and choose an MIC, or they may decide that the image is not readable.

```
[49]: df=UKMYC_PHENOTYPES.loc[(UKMYC_PHENOTYPES.PHENOTYPE_QUALITY=="LOW")]

len(df.loc[UKMYC_PHENOTYPES.BASHTHEBUGPRO_DILUTION.notna()])/len(df)
```

```
[49]: 0.7882357264925853
```

Note also, that the volunteers also finished looking at the images where PRIMARY\_DILUTION and AMYGDA\_DILUTION agree and hence these rows will have a BASHTHEBUG\_DILUTION reading even

though it does not affect the final DILUTION. It is provided mainly for Machine Learning from the images using the classifications as input features.

### 3.5 GROWTH

A useful by-product of reading all the images of plates with AMyGDA is that we measure the percentage of growth in the centre of each well on every plate we have an image for. This is stored here in a (long) table.

```
[50]: GROWTH=pandas.read_pickle(TABLES_PATH+"UKMYC_GROWTH.pkl.gz")
      print(len(GROWTH))
      GROWTH[:3]
```

2609184

```
[50]:
```

UNIQUEID	READINGDAY	DRUG	DILUTION	PLATEDESIGN \
site.01.subj.DR0013.lab.DR0013.iso.1	14	AMI	1	UKMYC6
			2	UKMYC6
			3	UKMYC6

UNIQUEID	READINGDAY	DRUG	DILUTION	SITEID \
site.01.subj.DR0013.lab.DR0013.iso.1	14	AMI	1	01
			2	01
			3	01


UNIQUEID	READINGDAY	DRUG	DILUTION	WELL_CONC \
site.01.subj.DR0013.lab.DR0013.iso.1	14	AMI	1	0.25
			2	0.50
			3	1.00

UNIQUEID	READINGDAY	DRUG	DILUTION	GROWTH
site.01.subj.DR0013.lab.DR0013.iso.1	14	AMI	1	9.554551
			2	0.270088
			3	0.129116

Let's quickly look at the distribution of measured growth in the control wells over time

```
[51]: GROWTH.reset_index(inplace=True)
      df=GROWTH.loc[(GROWTH.DRUG=="POS") & (GROWTH.DILUTION==0)]
      a=df['GROWTH'].hist(by=df.READINGDAY,figsize=(10,8),bins=50)
```



Note that there is a bias here; it is likely labs only allowed plates to incubate to 21 days if their growth at 14 days was poor so you cannot directly compare the two. We obviously could only show histograms for plates that were read at day 21 but that isn't what is shown here.

## 4 Genotype data

In this section I will quickly run through the tables containing the genetics information.

### 4.1 GENOMES

These contain one row per VCF file, i.e. one row per successful Clockwork output. Successful means the cluster process did not fail and all quality control checks were passed. Hence some samples have been excluded since they did not pass the overall genetic quality control metrics.

Remember several sets of short-reads could have been combined and processed together if `seq_rep` is something like `1_2_3`. This table is a join of what used to be two separate tables (`GENOMES` and `VCF_FILES`).

```
[52]: GENOMES=pandas.read_pickle(TABLES_PATH+"GENOMES.pkl.gz")
GENOMES[:3]
```

```
[52]:
```

	SITEID	SUBJID	LABID	ISOLATENO	\
UNIQUEID					
site.02.subj.0958.lab.22A197.iso.1	02	0958	22A197	1	
site.02.subj.0823.lab.2013241494.iso.1	02	0823	2013241494	1	
site.02.subj.0359.lab.222018-14.iso.1	02	0359	222018-14	1	

	SEQREPS	BELONGS_GPI	\
UNIQUEID			
site.02.subj.0958.lab.22A197.iso.1	197	True	
site.02.subj.0823.lab.2013241494.iso.1	241494	True	
site.02.subj.0359.lab.222018-14.iso.1	14222018	True	

	PER_SAMPLE_VCF_PRESENT	\
UNIQUEID		
site.02.subj.0958.lab.22A197.iso.1	True	
site.02.subj.0823.lab.2013241494.iso.1	True	
site.02.subj.0359.lab.222018-14.iso.1	True	

	REGENOTYPED_VCF_PRESENT	\
UNIQUEID		
site.02.subj.0958.lab.22A197.iso.1	True	
site.02.subj.0823.lab.2013241494.iso.1	True	
site.02.subj.0359.lab.222018-14.iso.1	True	

	CLOCKWORK_VERSION	TBI_INDEX	\
UNIQUEID			
site.02.subj.0958.lab.22A197.iso.1	0.8.3	True	
site.02.subj.0823.lab.2013241494.iso.1	0.8.3	True	
site.02.subj.0359.lab.222018-14.iso.1	0.8.3	True	

	KMER_COUNTS	SNP_DISTANCE_TO_H37rV	\
UNIQUEID			
site.02.subj.0958.lab.22A197.iso.1	False	1154.0	
site.02.subj.0823.lab.2013241494.iso.1	False	388.0	
site.02.subj.0359.lab.222018-14.iso.1	False	1147.0	

	SPECIES	LINEAGE_NAME	\
UNIQUEID			
site.02.subj.0958.lab.22A197.iso.1	M. tuberculosis	Lineage 2	
site.02.subj.0823.lab.2013241494.iso.1	M. tuberculosis	Lineage 4	
site.02.subj.0359.lab.222018-14.iso.1	M. tuberculosis	Lineage 2	

	SUBLINEAGE_NAME	LINEAGE_PERCENTAGE	\
UNIQUEID			



site.02.subj.0958.lab.22A197.iso.1	71.283784
site.02.subj.0823.lab.2013241494.iso.1	95.705521
site.02.subj.0359.lab.222018-14.iso.1	95.608108

	N_NULL	N_SNP	N_INDEL	N_FILTER_FAIL	\
UNIQUEID					
site.02.subj.0958.lab.22A197.iso.1	4934	0	1154		104
site.02.subj.0823.lab.2013241494.iso.1	2250	0	388		43
site.02.subj.0359.lab.222018-14.iso.1	3578	0	1147		118

	N_REF	N_HET	CATALOGUE_NAME	\
UNIQUEID				
site.02.subj.0958.lab.22A197.iso.1	0	14329	CRyPTIC	
site.02.subj.0823.lab.2013241494.iso.1	0	9442	CRyPTIC	
site.02.subj.0359.lab.222018-14.iso.1	0	13038	CRyPTIC	

	CATALOGUE_VERSION	TB_TYPE_1	\
UNIQUEID			
site.02.subj.0958.lab.22A197.iso.1		v1.31	MDR
site.02.subj.0823.lab.2013241494.iso.1		v1.31	UNK
site.02.subj.0359.lab.222018-14.iso.1		v1.31	UNK

	WGS_PREDICTION_STRING	\
UNIQUEID		
site.02.subj.0958.lab.22A197.iso.1	RRURRRSSSSSSSSSS	
site.02.subj.0823.lab.2013241494.iso.1	UUSUSSSSSSSSSSSS	
site.02.subj.0359.lab.222018-14.iso.1	SUSSSSSSUSSSSSSS	

IMAGE_MD5SUM	\
UNIQUEID	
site.02.subj.0958.lab.22A197.iso.1	{'02-0958-22A197-1-14': 'a587bac9ad2a0ebd36274...
site.02.subj.0823.lab.2013241494.iso.1	{'02-0823-2013241494-1-14': '698507bed7ff19268...
site.02.subj.0359.lab.222018-14.iso.1	{'02-0359-222018-14-1-14': '39c28529c7564ce379...

	FTP_PATH	\
UNIQUEID		
site.02.subj.0958.lab.22A197.iso.1	/well/bag/jeffk/release_staging/	
site.02.subj.0823.lab.2013241494.iso.1	/well/bag/jeffk/release_staging/	
site.02.subj.0359.lab.222018-14.iso.1	/well/bag/jeffk/release_staging/	

FTP_FILENAME_VCF	\
UNIQUEID	
site.02.subj.0958.lab.22A197.iso.1	00/01/41/00/14100/site.02.iso.1.subject.0958.l...

```

site.02.subj.0823.lab.2013241494.iso.1
00/01/41/43/14143/site.02.iso.1.subject.0823.l...
site.02.subj.0359.lab.222018-14.iso.1
00/01/08/73/10873/site.02.iso.1.subject.0359.l...

```

```

TREE_PATH \
UNIQUEID
site.02.subj.0958.lab.22A197.iso.1
dat/CRyPTIC2/V2/02/0958/22A197/1/regenotyped/
site.02.subj.0823.lab.2013241494.iso.1
dat/CRyPTIC2/V2/02/0823/2013241494/1/regenotyped/
site.02.subj.0359.lab.222018-14.iso.1
dat/CRyPTIC2/V2/02/0359/222018-14/1/regenotyped/

```

```

TREE_FILENAME_VCF \
UNIQUEID
site.02.subj.0958.lab.22A197.iso.1
site.02.subj.0958.lab.22A197.iso.1.v0.8.3.rege...
site.02.subj.0823.lab.2013241494.iso.1
site.02.subj.0823.lab.2013241494.iso.1.v0.8.3...
site.02.subj.0359.lab.222018-14.iso.1
site.02.subj.0359.lab.222018-14.iso.1.v0.8.3.r...

```

FASTQ\_MD5SUMS

```

UNIQUEID
site.02.subj.0958.lab.22A197.iso.1
site.02.subj.0823.lab.2013241494.iso.1
site.02.subj.0359.lab.222018-14.iso.1

```

```
[53]: len(GENOMES)
```

```
[53]: 70372
```

Since some of the later tables are VERY large, we also define a subset of 600 samples via this table. These were randomly chosen and so contain a mixture of sites as well as GPI/non-GPI etc.

```
[54]: GENOMES_SAMPLE=pandas.read_pickle('GENOMES_SAMPLE.pkl.gz')
len(GENOMES_SAMPLE)
```

```
[54]: 600
```

As for UKMYC\_PLATES, the hierarchical metadata fields SITEID, SUBJID, LABID, ISOLATENO, SEQREPS are included, the latter being specific to genetics.

The same Boolean flag, BELONGS\_GPI is also included here. This was used to decided which samples should be regenotyped (and which should not). Not all samples could be regenotyped due to the large memory requirements of the process. Martin Hunt is attempting to regenotype all 62k samples but this may not work.

If a sample was regenotyped, a ‘normal’ per-sample VCF file was also generated and hence samples with `BELONG_GPI==True` have 2 VCF files. The presence or absence of both types of files is indicated with `PER_SAMPLE_VCF_PRESENT` and `REGENOTYPED_VCF_PRESENT`.

If a sample `BELONGS_GPI` then both VCF files are processed in the sharded data tree and both have mini `VARIANTS` and `MUTATIONS` tables stored (thereby enabling potential comparisons between the two approaches). When constructing these tables, we take the approach of using the regenotyped VCF for the `BELONGS_GPI` samples and the per-sample VCF data for the remainder.

The exception is the 17 quality control samples from Comas and Gagneux; these have only been through the regenotyping process and hence have no per-sample VCF. These are also the reason why there are  $15228=15211+17$  samples with `BELONGS_GPI` in the `GENOMES` table. They can be identified by `SITEID=='QC'`.

Since the regenotyped VCF files are large, they often have an attendant `.tbi` index file. The presence of this is noted with `TBI_INDEX`. Likewise the presence of `kmer-counts.txt.gz` files are noted with `KMER_COUNTS`. These files have now been sorted as requested by Alex Lachapelle.

Again, like in `UKMYC_PLATES`, the path to the correct folder in the sharded tree is given by `TREE_PATH` and the filename of the VCF file is stored in `TREE_FILENAME_VCF`. Using the above Boolean flags, one can also construct paths to all the other files (e.g. kmer counts for machine learning).

Finally, the clockwork version is stored and also the md5sums for the `FASTQ` files are stored as JSON in `FASTQ_MD5SUMS`. This was necessary since there could be multiple pairs if there are multiple `SEQ_REPS` for this sample. For information, the path to the original vcf file provided by Jeff can be parsed from `(FTP_PATH,FTP_FILENAME_VCF)`.

```
[55]: pandas.crosstab(GENOMES.BELONGS_GPI,GENOMES.REGENOTYPED_VCF_PRESENT)
```

```
[55]: REGENOTYPED_VCF_PRESENT  False  True
BELONGS_GPI
False                        55144    0
True                         0  15228
```

```
[56]: pandas.crosstab(GENOMES.BELONGS_GPI,GENOMES.PER_SAMPLE_VCF_PRESENT)
```

```
[56]: PER_SAMPLE_VCF_PRESENT  False  True
BELONGS_GPI
False                        0  55144
True                        17  15211
```

This is the number of VCFs we have per site

```
[57]: GENOMES.SITEID.value_counts().sort_index()
```

```
[57]: 00      1615
      01      136
      02     1090
      03     1837
      04     4393
```

```

05      2930
06      2368
07     10435
08      1291
10      2662
11       463
13       286
14       389
16        69
17        91
20       454
21         7
ENA    39839
QC       17
Name: SITEID, dtype: int64

```

All the BELONGS\_GPI VCFs have a regenotyped VCF been passed through Sam Lipworth's SNPIT ([this version](#)) and hence this tables contains SPECIES, LINEAGE\_NAME and, where provided (mostly for Lineage 4), SUBLINEAGE\_NAME. Finally the LINEAGE\_PERCENTAGE is also given. If no regenotyped VCF is present, then these are all nulls.

Pleasingly, all samples belong to the *M. tuberculosis* complex!

```
[58]: GENOMES.loc[GENOMES.BELONGS_GPI].SPECIES.value_counts().sort_index()
```

```

[58]:
M. bovis BCG      0
M. bovis bovis    8
M. bovis caprae   3
M. bovis caprae   1
M. orygis         10
M. tuberculosis  15206
Name: SPECIES, dtype: int64

```

Of those predicted to be *M. tuberculosis*, the majority are Lineage 2 and 4, as expected.

```
[59]: GENOMES[GENOMES.SPECIES=="M. tuberculosis"].LINEAGE_NAME.value_counts().
      ↪sort_index()
```

```

[59]:
Lineage 1    1107
Lineage 2    5527
Lineage 3    1837
Lineage 4    6724
Lineage 5         2
Lineage 6         9
Name: LINEAGE_NAME, dtype: int64

```


Some sublineage information is available for Lineage 4

```
[60]: GENOMES[GENOMES.LINEAGE_NAME=="Lineage 4"].SUBLINEAGE_NAME.value_counts().
      ↪sort_values(ascending=False)
```

```
[60]: LAM          2235
      1554
      Haarlem     1344
      X-type       557
      S-type       524
      Tur          280
      Cameroon    118
      Ural         81
      Ghana        16
      Uganda       15
      Name: SUBLINEAGE_NAME, dtype: int64
```

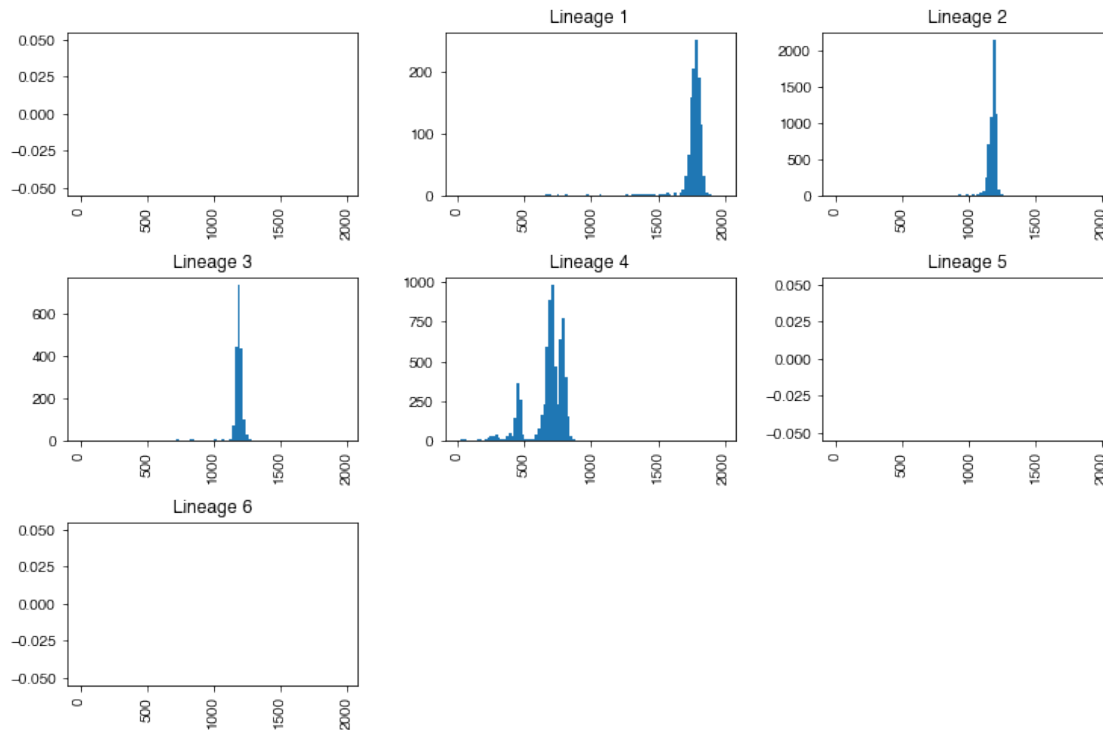
Again only for the BELONGS\_GPI samples which have regenotyped VCFs, the number of SNPs to the H37rV version 3 reference is recorded in SNP\_DISTANCE\_TO\_H37rV and we observe a number of peaks.

```
[61]: a=GENOMES[GENOMES.SPECIES=="M. tuberculosis"].SNP_DISTANCE_TO_H37rV.plot().
      ↪hist(bins=100)
```



Pleasingly, the different lineages explain the different peaks (although Lineages 2 & 3 overlap which just means they tend to be the same distance from the reference not that they are similar)

```
[62]: df=GENOMES[(GENOMES.SPECIES=="M. tuberculosis") & (GENOMES.LINEAGE_NAME.
↳isin(['Lineage 1','Lineage 2','Lineage 3','Lineage 4']))]
a=df['SNP_DISTANCE_TO_H37rV'].hist(by=df.
↳LINEAGE_NAME,bins=range(0,2000,20),figsize=(12,8))
```



Then we record some very high-level information about the antibiogram predicted from the default genetic catalogue.

```
[63]: GENOMES.CATALOGUE_NAME.value_counts()
```

```
[63]: CRyPTIC      70372
      Name: CATALOGUE_NAME, dtype: int64
```

```
[64]: GENOMES.CATALOGUE_VERSION.value_counts()
```

```
[64]: v1.31      70372
      Name: CATALOGUE_VERSION, dtype: int64
```

- TB\_TYPE\_1 provides a high-level description of the degree of resistance, based on the genetics. It shouldn't be relied upon but gives you an idea of the level of resistance.
- SUS susceptible according to the Bayesian approach in NEJM2018
- XDR resistant to RIF, INH, one of LEV or MXF and one of AMI or KAN
- MDR resistant to RIF, INH

- RIF resistant to RIF, susceptible to INH
- UNK everything else

```
[65]: GENOMES.loc[GENOMES.BELONGS_GPI].TB_TYPE_1.value_counts()
```

```
[65]: SUS      6665
      MDR      5429
      UNK      1896
      XDR       680
      RIF       558
      Name: TB_TYPE_1, dtype: int64
```

WGS\_PREDICTION\_STRING is simply the genetic antibiogram as a single string. The order of drugs is currently

```
["RIF", "INH", "PZA", "EMB", "AMI", "KAN", "LEV", "MXF", "ETH", "PAS", "RFB", "LZD", "BDQ", "DLM", "CFZ"]
```

Hence you can use slice the first four characters to compare to the NEJM paper. This is only supposed to be a hint and isn't definitive. For that please use the PREDICTIONS and EFFECTS tables which are described later.

## 4.2 VARIANTS

VARIANTS is a very long table since it contains all SNPs and INDELs detected by Clockwork across the whole genome for all samples, including those retrieved from the ENA. It contains calls parsed from both regenotyped and per-sample `vcf` files depending on what is available for that sample.

There is an approximate 1:1 mapping between a row in the VCF file making a call, and a row in this table. There are occasions where a row in the VCF can be parsed as containing e.g. two SNPs if `REF='cttcg'` and `ALT='ttttg'` and hence a single row in the VCF will map to two rows in the table.

It is not stored in `cryptic-tables/` since the gzipped CSV and PKL files are 2.8 GB and 1.6 GB respectively for the ~62k samples. We STRONGLY recommend you work with the PKL format since, unlike CSV, some of the fields are stored as categories which dramatically reduces the memory required. A table only containing the rows for the GPI samples is also included (`VARIANTS_GPI`: 18,159,378 rows : 326 MB and 623 MB).

Both tables are available on request, however for demonstration and testing purposes, the rows for the 600 randomly selected samples (1%) defined above in `GENOMES_SAMPLE` are stored in separate tables in `cryptic-tables/`. The whole table has currently (Aug 2020) 91,853,092 rows.

Since the regenotyped VCF files in particular yield large numbers of variants, we have adopted a 'mixed' approach; all SNPs and INDELs are recorded for all genes, however null calls and filter fails are only recorded for genes in the current resistance catalogue. To be clear a filter fail is a putative call that has not passed the quality and statistical thresholds set by Clockwork. The reason these are included for these genes is that one would like to know when predicting resistance e.g. one would treat a null or filter fail at position Ser450 in *rpoB* differently to a reference/wildtype call.

```
[66]: VARIANTS_SAMPLE=pandas.read_pickle(TABLES_PATH+"VARIANTS_SAMPLE.pkl.gz")
VARIANTS_SAMPLE[:3]
```

```
[66]:
```

UNIQUEID	VARIANT	IS_SNP	REF	ALT	GENOME_INDEX	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	True	c	g	1565	
	1977a>g	True	a	g	1977	
	4013t>c	True	t	c	4013	

UNIQUEID	VARIANT	GENE	ELEMENT_TYPE	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g			
	1977a>g	dnaN	GENE	
	4013t>c	recF	GENE	

UNIQUEID	VARIANT	MUTATION_TYPE	POSITION	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	SNP	NaN	
	1977a>g	SNP	-75.0	
	4013t>c	SNP	245.0	

UNIQUEID	VARIANT	NUCLEOTIDE_NUMBER	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	NaN	
	1977a>g	-75.0	
	4013t>c	734.0	

UNIQUEID	VARIANT	AMINO_ACID_NUMBER	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	NaN	
	1977a>g	NaN	
	4013t>c	245.0	

UNIQUEID	VARIANT	ASSOCIATED_WITH_GENE	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	False	
	1977a>g	True	
	4013t>c	True	

UNIQUEID	VARIANT	IN_PROMOTER	IN_CDS	IS_INDEL	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	False	False	False	
	1977a>g	True	False	False	
	4013t>c	False	True	False	

UNIQUEID	VARIANT	IS_HET	IS_NULL	\
----------	---------	--------	---------	---



site.02.subj.0345.lab.234051-15.iso.1	1565c>g	False	False
	1977a>g	False	False
	4013t>c	False	False

UNIQUEID	VARIANT	IS_FILTER_PASS	INDEL_LENGTH	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	True	NaN	
	1977a>g	True	NaN	
	4013t>c	True	NaN	

UNIQUEID	VARIANT	INDEL_1	INDEL_2	DP	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	None	None	238.0	
	1977a>g	None	None	240.0	
	4013t>c	None	None	212.0	

UNIQUEID	VARIANT	COVERAGE	DPF	FRS	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	238	1.1944	1.0	
	1977a>g	240	1.2044	1.0	
	4013t>c	212	1.0639	1.0	

UNIQUEID	VARIANT	GT_CONF	\
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	1870.410034	
	1977a>g	1882.469971	
	4013t>c	1712.030029	

UNIQUEID	VARIANT	GT_CONF_PERCENTILE	SITEID
site.02.subj.0345.lab.234051-15.iso.1	1565c>g	88.739998	02
	1977a>g	89.720001	02
	4013t>c	67.720001	02

```
[67]: len(VARIANTS_SAMPLE)
```

```
[67]: 676942
```

To help with some of the graphs below, let's join to the `GENOMES_SAMPLE` table so we can add the `BELONGS_GPI` column.

```
[68]: def assign_gpi_description(row):
        if row['BELONGS_GPI']:
            return("GPI")
        else:
            return("NOT GPI")
```

```

GENOMES_SAMPLE['GPI_LABEL']=GENOMES_SAMPLE.apply(assign_gpi_description,axis=1)
GENOMES_SAMPLE[:3]

VARIANTS_SAMPLE.reset_index(inplace=True)
VARIANTS_SAMPLE.set_index('UNIQUEID',inplace=True)
VARIANTS_SAMPLE=VARIANTS_SAMPLE.join(GENOMES_SAMPLE[['GPI_LABEL']],how='left')
VARIANTS_SAMPLE.GPI_LABEL.value_counts()

```

```

[68]: NOT GPI      515924
      GPI         161018
      Name: GPI_LABEL, dtype: int64

```

Genetic variants are by definition all located on the reference H37rV genome (version 3) using GENOME\_INDEX and therefore this *always* contains a value.

```

[69]: VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.GENOME_INDEX.isna()]

```

```

[69]: Empty DataFrame
Columns: [VARIANT, IS_SNP, REF, ALT, GENOME_INDEX, GENE, ELEMENT_TYPE,
MUTATION_TYPE, POSITION, NUCLEOTIDE_NUMBER, AMINO_ACID_NUMBER,
ASSOCIATED_WITH_GENE, IN_PROMOTER, IN_CDS, IS_INDEL, IS_HET, IS_NULL,
IS_FILTER_PASS, INDEL_LENGTH, INDEL_1, INDEL_2, DP, COVERAGE, DPF, FRS, GT_CONF,
GT_CONF_PERCENTILE, SITEID, GPI_LABEL]
Index: []

```

A genetic ‘variant’ is either (i) a single nucleotide polymorphism or (ii) an insertion or deletion of a specified number of nucleotides or (iii) a null (which could be either a SNP or an INDEL or nothing, we don’t know!). These can be identified via MUTATION\_TYPE and also (redundantly) using the Booleans IS\_SNP and IS\_INDEL.

```

[70]: pandas.crosstab(VARIANTS_SAMPLE.MUTATION_TYPE,VARIANTS_SAMPLE.IS_SNP)

```

```

[70]: IS_SNP      False   True
      MUTATION_TYPE
      INDEL      57814      0
      NULL      17385      0
      SNP         0  601743

```

```

[71]: pandas.crosstab(VARIANTS_SAMPLE.MUTATION_TYPE,VARIANTS_SAMPLE.IS_INDEL)

```

```

[71]: IS_INDEL      False   True
      MUTATION_TYPE
      INDEL         0  57814
      NULL      17385      0
      SNP      601743      0

```

```

[72]: pandas.crosstab(VARIANTS_SAMPLE.MUTATION_TYPE,VARIANTS_SAMPLE.IS_NULL)

```

```
[72]: IS_NULL      False  True
      MUTATION_TYPE
      INDEL      57814    0
      NULL       0    17385
      SNP       601743    0
```

SNPs simply have the nucleotide of the reference genome in REF and the observed allele in ALT.

Note that the allowed REF bases are [a,t,c,g] but the allowed ALT bases are [a,t,c,g,o,x,z] where o indicates a vcf filter fail, x indicates a Null call and z a Heterogenous call. As we shall see later, Clockwork is not, at present, making any Het calls and therefore there are no zs but the code allows for them. There are the associated IS\_FILTER\_PASS,IS\_NULL and IS\_HET Boolean fields to help with identifying such variants.

```
[73]: df=VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.IS_SNP]
      pandas.crosstab(df.REF,df.ALT)
```

```
[73]: ALT      a      c      g      o      t
      REF
      a      0  26303  101066  225    5557
      c    35527      0   38980  274   100215
      g    87828  48653      0   326    33239
      t     3385  91772   28228  165      0
```

SNPs are simply described as GENOME\_INDEX REF>ALT in the VARIANT field e.g. 1849c>a whilst indels are simply noted e.g. 1849\_indel.

(UNIQUEID,VARIANT) is therefore the (unique) primary key.

```
[74]: VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.IS_INDEL][:3]
```

```
[74]:
```

	VARIANT	IS_SNP	REF	ALT	\
UNIQUEID					
site.00.subj.LE10KTB_23.lab.7627572.iso.1	26747_indel	False	gc	g	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	34568_indel	False	tc	t	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	49690_indel	False	gcc	g	

	GENOME_INDEX	GENE	ELEMENT_TYPE	\
UNIQUEID				
site.00.subj.LE10KTB_23.lab.7627572.iso.1	26747	Rv0021c	LOCUS	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	34568	bioF2	GENE	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	49690	Rv0045c	LOCUS	

	MUTATION_TYPE	POSITION	\
UNIQUEID			
site.00.subj.LE10KTB_23.lab.7627572.iso.1	INDEL	135.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	INDEL	274.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	INDEL	250.0	

UNIQUEID	NUCLEOTIDE_NUMBER	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	135.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	274.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	250.0	

UNIQUEID	AMINO_ACID_NUMBER	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	45.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	92.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	84.0	

UNIQUEID	ASSOCIATED_WITH_GENE	IN_PROMOTER	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	

UNIQUEID	IN_CDS	IS_INDEL	IS_HET	IS_NULL	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	True	False	False	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	True	False	False	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	True	False	False	

UNIQUEID	IS_FILTER_PASS	INDEL_LENGTH	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	-1.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	-1.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	-2.0	

UNIQUEID	INDEL_1	INDEL_2	DP	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	26747_del	26747_del_1	26.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	34568_del	34568_del_1	22.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	49690_del	49690_del_2	19.0	

UNIQUEID	COVERAGE	DPF	FRS	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	23	1.1685	0.9231	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	22	0.9887	1.0000	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	19	0.8539	1.0000	

UNIQUEID	GT_CONF	GT_CONF_PERCENTILE	\
site.00.subj.LE10KTB_23.lab.7627572.iso.1	153.789993	43.340000	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	178.750000	56.919998	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	159.929993	46.880001	

UNIQUEID	SITEID	GPI_LABEL
site.00.subj.LE10KTB_23.lab.7627572.iso.1	00	NOT GPI
site.00.subj.LE10KTB_23.lab.7627572.iso.1	00	NOT GPI
site.00.subj.LE10KTB_23.lab.7627572.iso.1	00	NOT GPI

Whilst INDELs have one or more nucleotides from the reference in REF and one or more nucleotides in ALT and the first position is assumed to be the start of the INDEL, which may not be true, but is the most straightforward assumption and otherwise you get tangled up in dividing variants into one or more constituents (since INDEL and SNP are not an orthogonal basis set).

The length of the INDEL is also recorded; note that this is the *net* length i.e. `len(ALT)-len(REF)`

```
[75]: VARIANTS_SAMPLE.INDEL_LENGTH.value_counts().sort_index()
```

```
[75]: -29750.0    1
      -12719.0    5
      -7890.0     1
      -6967.0     1
      -6807.0     1
      ..
       6166.0     1
       8335.0     1
      11044.0     1
      13044.0     1
      14054.0     1
      Name: INDEL_LENGTH, Length: 443, dtype: int64
```

Since the nomenclature for an INDEL forms a nature hierarchy and we've used the simplest descriptor in the VARIANT field, the descending levels are included in INDEL\_1 and INDEL\_2.

```
[76]: VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.
      ↪IS_INDEL][['VARIANT', 'INDEL_1', 'INDEL_2', 'INDEL_LENGTH']][:3]
```

```
[76]:
```

	VARIANT	INDEL_1	\
UNIQUEID			
site.00.subj.LE10KTB_23.lab.7627572.iso.1	26747_indel	26747_del	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	34568_indel	34568_del	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	49690_indel	49690_del	

	INDEL_2	INDEL_LENGTH
UNIQUEID		
site.00.subj.LE10KTB_23.lab.7627572.iso.1	26747_del_1	-1.0
site.00.subj.LE10KTB_23.lab.7627572.iso.1	34568_del_1	-1.0
site.00.subj.LE10KTB_23.lab.7627572.iso.1	49690_del_2	-2.0

Each successive descriptor gives a little more information; first about whether it is an insertion or a deletion, and then how many bases are involved.

At present, there are no examples (beyond frameshifts) where different ‘flavours’ of INDELs at the same `GENOME_POSITION` need to be distinguished since they have been associated with different effects on drug resistance. This is likely to change (or at least be tested), hence the flexibility is built in here.

All SNP and INDEL variants may (or may not) fall within a coding region (gene) or its promoter. This can be identified by the Boolean `ASSOCIATED_WITH_GENE`. If `True`, then `ELEMENT_TYPE` can be used to distinguish the ‘type’ of coding region. Note that being in the ‘promoter’ is ill-defined and here is assumed to be 100 bases upstream of the start codon (or up to the next coding region, whichever comes sooner).

```
[77]: VARIANTS_SAMPLE.ASSOCIATED_WITH_GENE.value_counts()
```

```
[77]: True      620750
      False    56192
      Name: ASSOCIATED_WITH_GENE, dtype: int64
```

```
[78]: pandas.crosstab(VARIANTS_SAMPLE.ELEMENT_TYPE,VARIANTS_SAMPLE.
      ↪ASSOCIATED_WITH_GENE)
```

```
[78]: ASSOCIATED_WITH_GENE  False   True
      ELEMENT_TYPE
                        56192      0
      GENE                0  339352
      LOCUS                0  279085
      RNA                  0   2313
```

GENE, LOCUS and RNA are as defined in the H37rV Genbank file (NC\_000962.3.gbk). The RNA genes found are, of course, ribosomal. Implicit, therefore, is the assumption that GENE and LOCUS code proteins, whilst RNA do not.

GENE encodes the name of the GENE or LOCUS as defined by the H37rV Genbank file. There are examples of *M. tuberculosis* genes that are referred to in the literature by one name, but are called something else in the GenBank file. Only using the latter makes any sense!

```
[79]: VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.ELEMENT_TYPE=='RNA'].GENE.value_counts()
```

```
[79]: rrl      1161
      rrs      1149
      rrf        3
      zwf2       0
      Rv2235      0
      ...
      fadD19      0
      fadD18      0
      fadD17      0
      fadD15      0
      0
      Name: GENE, Length: 3863, dtype: int64
```

If we only considering variants that are associated with a gene/locus, we now find that there are multiple ways of identifying the genetic position where the variant occurs; the position in the whole genome (`GENOME_INDEX`) but also the number of nucleotides since the start of the gene/locus (`NUCLEOTIDE_NUMBER` – this is negative by definition for promoters) and, if the gene/locus encodes protein, the amino acid number (`AMINO_ACID_NUMBER`).

Note that this also makes it clear why one must use the *M. tuberculosis* (i.e. reference) numbering, and why e.g. using *E. coli* numbering for *rpoB* is confusing and prevents matching to the GenBank reference.

Note also that if these cannot be defined for a variant (e.g. it isn't associated with a gene, or encodes RNA and therefore `AMINO_ACID_NUMBER` is nonsensical) then you'll find a `NaN`. It is a peculiarity of Pandas that only floats can hold NaNs, whilst ints cannot, and therefore these are all stored as floats.

- `NUCLEOTIDE_NUMBER` is simply the 1-based number of the base. A peculiarity of genes is there is no 0, so if there is a promoter, it will run -3,-2,-1,1,2,3,4...
- `AMINO_ACID_NUMBER` is the sequential number of the amino acid residue that the base belongs to/codes for. Hence it is only populated in the coding region of genes that code for protein (i.e. not RNA encoding genes like *rrs*). For the same sequence as above it will be, `NaN,NaN,NaN,1,1,1,2`

To facilitate joining between the `VARIANT` and `MUTATIONS` tables there is fourth aggregated location field called `POSITION`. If the variant occurs in the coding region of a gene that codes protein this is `AMINO_ACID_NUMBER`, otherwise it is simply `NUCLEOTIDE_NUMBER` and if the variant is not associated with a gene in any way, then it is a `NaN`.

```
[80]: VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.
      ↪ASSOCIATED_WITH_GENE][["GENOME_INDEX", "NUCLEOTIDE_NUMBER", "AMINO_ACID_NUMBER", "POSITION"]][
      ↪5]
```

```
[80]:
```

	GENOME_INDEX	NUCLEOTIDE_NUMBER	\
UNIQUEID			
site.00.subj.LE10KTB_23.lab.7627572.iso.1	1977	-75.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	3446	167.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	4013	734.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	7362	61.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	7585	284.0	

	AMINO_ACID_NUMBER	POSITION
UNIQUEID		
site.00.subj.LE10KTB_23.lab.7627572.iso.1	NaN	-75.0
site.00.subj.LE10KTB_23.lab.7627572.iso.1	56.0	56.0
site.00.subj.LE10KTB_23.lab.7627572.iso.1	245.0	245.0
site.00.subj.LE10KTB_23.lab.7627572.iso.1	21.0	21.0
site.00.subj.LE10KTB_23.lab.7627572.iso.1	95.0	95.0

Two other Booleans are provided to help you; `IN_CDS` and `IN_PROMOTER` which do what you think they do!

Finally, each variant has stored the quality metrics used by Clockwork in deciding whether or not to make a call. I will describe these, but since I am not an expert at the meanings of these, please direct questions to Zam, Jeff or Martin.

The total depth from the pile-up is stored in DP whilst the aggregate depth of the top two alleles is stored in COVERAGE. Most of the time these are identical, but on occasion DP>COVERAGE presumably because more than two bases were observed in pile-up.

```
[81]: VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.DP>VARIANTS_SAMPLE.COVERAGE][:3]
```

```
[81]:
```

	VARIANT	IS_SNP	REF	ALT	\
UNIQUEID					
site.00.subj.LE10KTB_23.lab.7627572.iso.1	69984c>a	True	c	a	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	69989g>a	True	g	a	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	72549g>a	True	g	a	

	GENOME_INDEX	GENE	ELEMENT_TYPE	\
UNIQUEID				
site.00.subj.LE10KTB_23.lab.7627572.iso.1	69984	Rv0064	LOCUS	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	69989	Rv0064	LOCUS	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	72549	icd2	GENE	

	MUTATION_TYPE	POSITION	\
UNIQUEID			
site.00.subj.LE10KTB_23.lab.7627572.iso.1	SNP	455.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	SNP	457.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	SNP	655.0	

	NUCLEOTIDE_NUMBER	\
UNIQUEID		
site.00.subj.LE10KTB_23.lab.7627572.iso.1	1365.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	1370.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	1963.0	

	AMINO_ACID_NUMBER	\
UNIQUEID		
site.00.subj.LE10KTB_23.lab.7627572.iso.1	455.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	457.0	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	655.0	

	ASSOCIATED_WITH_GENE	IN_PROMOTER	\
UNIQUEID			
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	

	IN_CDS	IS_INDEL	IS_HET	IS_NULL	\
UNIQUEID					



site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	False	False
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	False	False
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	False	False	False


	IS_FILTER_PASS	INDEL_LENGTH	\
UNIQUEID			
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	NaN	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	NaN	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	True	NaN	

	INDEL_1	INDEL_2	DP	COVERAGE	\
UNIQUEID					
site.00.subj.LE10KTB_23.lab.7627572.iso.1	None	None	32.0	30	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	None	None	32.0	30	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	None	None	25.0	24	

	DPF	FRS	GT_CONF	\
UNIQUEID				
site.00.subj.LE10KTB_23.lab.7627572.iso.1	1.4381	1.00	220.899994	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	1.4381	1.00	220.899994	
site.00.subj.LE10KTB_23.lab.7627572.iso.1	1.1235	0.96	181.759995	


	GT_CONF_PERCENTILE	SITEID	GPI_LABEL
UNIQUEID			
site.00.subj.LE10KTB_23.lab.7627572.iso.1	76.120003	00	NOT GPI
site.00.subj.LE10KTB_23.lab.7627572.iso.1	76.120003	00	NOT GPI
site.00.subj.LE10KTB_23.lab.7627572.iso.1	58.619999	00	NOT GPI

```
[82]: a=VARIANTS_SAMPLE['DP'].hist(bins=np.arange(0,500,10),figsize=(6,4))
```



The depth at that position as a fraction of the average depth over the whole genome is stored in DPF. This is therefore a float. As you'd expect this is centred on unity, but there is also a peak centred on zero that are all null calls.

```
[83]: a=VARIANTS_SAMPLE['DPF'].hist(bins=numpy.arange(0,2,0.05),figsize=(6,4))
```




```
[84]: a=VARIANTS_SAMPLE.loc[VARIANTS_SAMPLE.DPF<0.01]
pandas.crosstab(a.IS_NULL,a.IS_FILTER_PASS)
```

```
[84]: IS_FILTER_PASS  False  True
IS_NULL
False              23    413
True             1118  13533
```

Then we have the Fraction of Read Support (FRS) which by definition has an upper bound of unity. This is the closest Clockwork gets to thinking about het calls at present.


```
[85]: # a=VARIANTS_SAMPLE.FRS.hist(bins=numpy.arange(0,1.1,0.02))
a=VARIANTS_SAMPLE['FRS'].hist(bins=numpy.arange(0.0,1.1,0.
↪02),figsize=(6,4),log=True)
```



Note that everything with  $FRS < 0.9$  is either a NULL or FILTER\_FAIL call.

Internally, Clockwork uses a model to predict the confidence of the call; this is stored in GT\_CONF. Because the GPI samples are called together (by definition!) we have to start treating the GPI/NOT GPI calls separately from now on and as the quantities and thresholds are NOT equivalent.

```
[87]: a=VARIANTS_SAMPLE.GT_CONF.hist(by=VARIANTS_SAMPLE.
↪GPI_LABEL,bins=range(0,5000,100),figsize=(12,4))
```




Since the threshold one might use to discard low confidence values is itself a function of depth, Clockwork creates a set of reads with the same average depth as the sample and then analyses these to create a `GT_CONF` distribution. It then uses this to convert the `GT_CONF` values for the actual sample into percentiles, which are stored as `GT_CONF_PERCENTILE` and the threshold is set as a percentile value.

Unfortunately `GT_CONF_PERCENTILE` cannot be compared between regenotyped and per-sample `vcf` files. For the latter a threshold of 5% was applied and everything below that was recorded as a filter fail. For the former a smaller threshold was applied by considering the precision/recall of the 17 high-quality QC Comas samples.

Hence `GT_CONF_PERCENTILE` is the preferred metric since it accounts for the depth in each sample and runs 0-100.


```
[88]: a=VARIANTS_SAMPLE.GT_CONF_PERCENTILE.hist(by=VARIANTS_SAMPLE.
      ↪GPI_LABEL,bins=range(0,100,1),figsize=(12,4))
```



There are more nulls in the GPI simply because during the regenotyping each sample is examined at any and every position where there is evidence of a call in any sample, hence in many samples either a reference (wildtype) or null call is returned depending on the pile-up.

The GT\_CONF\_PERCENTILE thresholds for definite calls are 0.25% for the GPI and 5.0% for the non-GPI.

```
[89]: df=VARIANTS_SAMPLE.loc[(VARIANTS_SAMPLE.IS_FILTER_PASS) & (~VARIANTS_SAMPLE.  
    ↪ IS_NULL)]  
a=df.GT_CONF_PERCENTILE.hist(by=df.GPI_LABEL,bins=numpy.arange(0,10,0.  
    ↪ 1),figsize=(12,4))
```



### 4.3 MUTATIONS

Although not as large as VARIANTS, the MUTATIONS table is still large at 82,290,716 rows (Aug 2020) and takes minutes to load on my workstation with 48 GB of memory. The compressed PKL and CSV versions take up 736 MB and 1.33 GB, respectively on disc. Again we STRONGLY recommend you work with the PKL version if you can to avoid memory issues. Like VARIANTS the GPI subset is stored. This is small enough when compressed (15,977,222 rows: 124 MB / 275 MB) that it is stored in cryptic-tables/.

Finally, the rows relating to the 1% sample defined above in GENOMES\_SAMPLE are also stored for testing and demonstration in MUTATIONS\_SAMPLE.

```
[90]: MUTATIONS=pandas.read_pickle(TABLES_PATH+"MUTATIONS_SAMPLE.pkl.gz")  
MUTATIONS[:3]
```

```
[90]:
```

UNIQUEID	GENE	MUTATION	POSITION \
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	31.0
	PE1	L485L	485.0
	PE3	T14A	14.0

UNIQUEID	GENE	MUTATION	AMINO_ACID_NUMBER \
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	31.0
	PE1	L485L	485.0

	PE3	T14A		14.0
--	-----	------	--	------

			GENOME_INDEX	\
UNIQUEID	GENE	MUTATION		
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	NaN	
	PE1	L485L	NaN	
	PE3	T14A	NaN	

			NUCLEOTIDE_NUMBER	\
UNIQUEID	GENE	MUTATION		
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	NaN	
	PE1	L485L	NaN	
	PE3	T14A	NaN	

			REF	ALT	IS_SNP	\
UNIQUEID	GENE	MUTATION				
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	caa	cga	True	
	PE1	L485L	ctg	ttg	True	
	PE3	T14A	acg	gcg	True	

			IS_INDEL	IN_CDS	\	
UNIQUEID	GENE	MUTATION				
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	False	True		
	PE1	L485L	False	True		
	PE3	T14A	False	True		

			IN_PROMOTER	\		
UNIQUEID	GENE	MUTATION				
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	False			
	PE1	L485L	False			
	PE3	T14A	False			

			IS_SYNONYMOUS	\		
UNIQUEID	GENE	MUTATION				
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	False			
	PE1	L485L	True			
	PE3	T14A	False			

			IS_NONSYNONYMOUS	\		
UNIQUEID	GENE	MUTATION				
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	True			
	PE1	L485L	False			
	PE3	T14A	True			

			IS_HET	IS_NULL	\	
UNIQUEID	GENE	MUTATION				
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	False	False		

	PE1	L485L	False	False
	PE3	T14A	False	False

			IS_FILTER_PASS	\
UNIQUEID	GENE	MUTATION		
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	True	
	PE1	L485L	True	
	PE3	T14A	True	

			ELEMENT_TYPE	\
UNIQUEID	GENE	MUTATION		
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	GENE	
	PE1	L485L	GENE	
	PE3	T14A	GENE	

			MUTATION_TYPE	\
UNIQUEID	GENE	MUTATION		
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R	AAM	
	PE1	L485L	AAM	
	PE3	T14A	AAM	

			INDEL_LENGTH	INDEL_1	\
UNIQUEID	GENE	MUTATION			
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R		NaN	
	PE1	L485L		NaN	
	PE3	T14A		NaN	

			INDEL_2	SITEID	\
UNIQUEID	GENE	MUTATION			
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R		02	
	PE1	L485L		02	
	PE3	T14A		02	

NUMBER_NUCLEOTIDE_CHANGES				
UNIQUEID	GENE	MUTATION		
site.02.subj.0345.lab.234051-15.iso.1	35kd_ag	Q31R		
1				
	PE1	L485L		
1				
	PE3	T14A		
1				

The primary key (i.e. unique) is `UNIQUEID,GENE,MUTATION`. This protein level view (i.e. amino acids) needs to be separate from `VARIANTS` since you can have up to three SNPs in a single codon (and therefore three rows in `VARIANTS`) which would be represented by a single row here in `MUTATIONS`. This does mean, however, that if you want to find out the (min or max) `COVERAGE` in a codon, you need to join back to `VARIANTS`. This makes clear that the quality information only makes sense at

the nucleotide level, not the codon level.

For more information on the grammar used to describe each mutation, head [here](#).

Let's look at the mutations in the *rpoB* RRDR and look for any mutations where more than one base in the codon are different compared to the reference.

```
[91]: MUTATIONS.reset_index(inplace=True)
df=MUTATIONS.loc[(MUTATIONS.GENE=='rpoB') & (MUTATIONS.POSITION>=428) &
↳ (MUTATIONS.POSITION<=452) & (~MUTATIONS.IS_NULL)]
pandas.crosstab(df.MUTATION,df.NUMBER_NUCLEOTIDE_CHANGES,margins=True)
```

```
[91]: NUMBER_NUCLEOTIDE_CHANGES    1   2  All
MUTATION
D435A                1   0   1
D435V               11   0  11
H445D                9   0   9
H445L                1   0   1
H445S                0   2   2
H445Y                7   0   7
L430P                3   0   3
L452P                5   0   5
Q429H                1   0   1
Q432P                1   0   1
S450F                0   2   2
S450L               83   0  83
S450O                4   0   4
S450W                1   0   1
All                127   4  131
```

All the mutations in our sample are SNPs in *rpoB* with 83 out of 131 being S450L, as expected and only four amino acid mutations involve more than 1 nucleotide change in the codon compared to reference.

Most of the remaining fields are there to help you select the mutations you want. These include a series of Booleans that are hopefully obvious: IS\_SNP, IS\_INDEL, IN\_CDS, IN\_PROMOTER, IS\_SYNONYMOUS, IS\_NONSYNOYMOUS, IS\_HET, IS\_NULL. And, yes, many of these are redundant e.g. IN\_CDS=~IN\_PROMOTER but it just makes life a bit easier.

ELEMENT\_TYPE is the same as in VARIANTS, as are the additional descriptors for INDELS: INDEL\_LENGTH, INDEL\_1 and INDEL\_2.

MUTATION\_TYPE is new and distinguishes between an amino acid mutation (AAM) e.g. a codon change which may or may not be synonymous and a nucleotide SNP e.g. in a promoter or an RNA gene, as well as INDELS, which can be anywhere.

```
[92]: MUTATIONS.MUTATION_TYPE.value_counts()
```

```
[92]: AAM      526701
      INDEL    45175
```



```
SNP          38361
Name: MUTATION_TYPE, dtype: int64
```

Note that since FILTER\_FAIL and NULL calls are included for genes in the resistance catalogue these need to be distinguished. This is done via the ‘artificial’ amino acids O and X, respectively. At present a NULL (i.e. x) in a codon trumps a FILTER\_FAIL (i.e. o) . Hence an ALT of aox is translated to a X i.e. a NULL. There are only 8 instances in a our sample, however.

```
[94]: MUTATIONS.loc[(MUTATIONS.ALT.str.contains('x')) & (MUTATIONS.ALT.str.
↳contains('o'))]
```

```
[94]:
```

	UNIQUEID	GENE	MUTATION	POSITION	\
16247	site.05.subj.CA-1366.lab.CO-02570-19.iso.1	pepQ	A90X	90.0	
79056	site.04.subj.01017.lab.720678.iso.1	embA	G149X	149.0	
79432	site.04.subj.01017.lab.720678.iso.1	ethR	W145X	145.0	
80045	site.04.subj.01017.lab.720678.iso.1	mmpL3	M211X	211.0	
80085	site.04.subj.01017.lab.720678.iso.1	mmpL3	P867X	867.0	
80199	site.04.subj.01017.lab.720678.iso.1	mshA	L277X	277.0	
547769	site.03.subj.DR-8.lab.IML-01060.iso.1	ethA	V191X	191.0	

	AMINO_ACID_NUMBER	GENOME_INDEX	NUCLEOTIDE_NUMBER	REF	ALT	IS_SNP	\
16247	90.0	NaN	NaN	gcc	xoc	True	
79056	149.0	NaN	NaN	ggc	oxx	True	
79432	145.0	NaN	NaN	tgg	oxg	True	
80045	211.0	NaN	NaN	atg	otx	True	
80085	867.0	NaN	NaN	ccg	xog	True	
80199	277.0	NaN	NaN	ctg	cox	True	
547769	191.0	NaN	NaN	gtg	xto	True	

	IS_INDEL	IN_CDS	IN_PROMOTER	IS_SYNONYMOUS	IS_NONSYNONYMOUS	\
16247	False	True	False	False	True	
79056	False	True	False	False	True	
79432	False	True	False	False	True	
80045	False	True	False	False	True	
80085	False	True	False	False	True	
80199	False	True	False	False	True	
547769	False	True	False	False	True	

	IS_HET	IS_NULL	IS_FILTER_PASS	ELEMENT_TYPE	MUTATION_TYPE	\
16247	False	True	True	GENE	AAM	
79056	False	True	True	GENE	AAM	
79432	False	True	True	GENE	AAM	
80045	False	True	True	GENE	AAM	
80085	False	True	True	GENE	AAM	
80199	False	True	True	GENE	AAM	
547769	False	True	True	GENE	AAM	

	INDEL_LENGTH	INDEL_1	INDEL_2	SITEID	NUMBER_NUCLEOTIDE_CHANGES
16247	NaN			05	2
79056	NaN			04	3
79432	NaN			04	2
80045	NaN			04	2
80085	NaN			04	2
80199	NaN			04	2
547769	NaN			03	2

#### 4.4 EFFECTS

Now that we have a comprehensive view of all the genetic variants and their associated protein amino acid changes, we can apply one or more genetic resistance catalogues.

We are currently applying the CRyPTICv1.31 catalogue which is a merged catalogue comprising NEJM2018 for the first-line compounds and ERJ2017 for the rest. It also includes all the genes identified by the Seq&Treat project as being of potential interest. As mentioned above all of these genes therefore have nulls and filter fails recorded in VARIANTS and MUTATIONS although since they only have default rows in the catalogue they can only ever cause a U or S to be returned.

```
[95]: EFFECTS=pandas.read_pickle(TABLES_PATH+"EFFECTS.pkl.gz")
      EFFECTS[:10]
```

```
[95]:
```

UNIQUEID	CATALOGUE_VERSION	CATALOGUE_GRAMMAR	SITEID \	DRUG	GENE	MUTATION	CATALOGUE_NAME
site.02.subj.0958.lab.22A197.iso.1	PZA	PPE35	A3910	CRyPTIC	v1.31		
GARC1	02		L896S	CRyPTIC	v1.31		
GARC1	02						
GARC1	02	DLM	Rv1816	D70G	CRyPTIC	v1.31	
GARC1	02	BDQ	Rv1816	D70G	CRyPTIC	v1.31	
GARC1	02	CFZ	Rv1816	D70G	CRyPTIC	v1.31	
GARC1	02	PAS	Rv1816	D70G	CRyPTIC	v1.31	
GARC1	02	LZD	Rv1816	D70G	CRyPTIC	v1.31	
GARC1	02	PZA	Rv3236c	T102A	CRyPTIC	v1.31	
GARC1	02	AMI	aftB	D397G	CRyPTIC	v1.31	
GARC1	02	KAN	aftB	D397G	CRyPTIC	v1.31	
GARC1	02						

UNIQUEID	PREDICTION	DRUG	GENE	MUTATION	CATALOGUE_NAME
CATALOGUE_VERSION	CATALOGUE_GRAMMAR				
site.02.subj.0958.lab.22A197.iso.1	PZA	PPE35	A3910	CRyPTIC	v1.31
GARC1	U		L896S	CRyPTIC	v1.31
GARC1	U				
		DLM	Rv1816	D70G	CRyPTIC
GARC1	U				v1.31
		BDQ	Rv1816	D70G	CRyPTIC
GARC1	U				v1.31
		CFZ	Rv1816	D70G	CRyPTIC
GARC1	U				v1.31
		PAS	Rv1816	D70G	CRyPTIC
GARC1	U				v1.31
		LZD	Rv1816	D70G	CRyPTIC
GARC1	U				v1.31
		PZA	Rv3236c	T102A	CRyPTIC
GARC1	U				v1.31
		AMI	aftB	D397G	CRyPTIC
GARC1	U				v1.31
		KAN	aftB	D397G	CRyPTIC
GARC1	U				v1.31

EFFECTS contains one row per mutation in each catalogue gene per associated drug for a defined version of a single catalogue. Hence a single row in MUTATIONS e.g. *gyrA*\_A90V may result in multiple rows in EFFECTS since not only can that mutation be associated with resistance to several fluoroquinolones but also a range of different catalogues, perhaps also different versions of a single catalogue, may have been applied. In addition, there may be other *gyrA* mutations in the same sample, each of which will contribute one (or more) row to EFFECTS. Consider this sample which has 4 mutations in *gyrA*.

```
[96]: EFFECTS.reset_index(inplace=True)
      EFFECTS.loc[(EFFECTS.GENE=='gyrA') & (EFFECTS.UNIQUEID=="site.02.subj.0914.lab.
      ↪22A148.iso.1")]
```

```
[96]:
```

	UNIQUEID	DRUG	GENE	MUTATION	CATALOGUE_NAME	\
11701	site.02.subj.0914.lab.22A148.iso.1	MXF	gyrA	E21Q	CRyPTIC	
11702	site.02.subj.0914.lab.22A148.iso.1	OFX	gyrA	E21Q	CRyPTIC	
11703	site.02.subj.0914.lab.22A148.iso.1	LEV	gyrA	E21Q	CRyPTIC	
11704	site.02.subj.0914.lab.22A148.iso.1	MXF	gyrA	A90V	CRyPTIC	
11705	site.02.subj.0914.lab.22A148.iso.1	OFX	gyrA	A90V	CRyPTIC	
11706	site.02.subj.0914.lab.22A148.iso.1	LEV	gyrA	A90V	CRyPTIC	
11707	site.02.subj.0914.lab.22A148.iso.1	MXF	gyrA	S95T	CRyPTIC	
11708	site.02.subj.0914.lab.22A148.iso.1	OFX	gyrA	S95T	CRyPTIC	
11709	site.02.subj.0914.lab.22A148.iso.1	LEV	gyrA	S95T	CRyPTIC	
11710	site.02.subj.0914.lab.22A148.iso.1	MXF	gyrA	G668D	CRyPTIC	

11711	site.02.subj.0914.lab.22A148.iso.1	OFX	gyrA	G668D	CRyPTIC
11712	site.02.subj.0914.lab.22A148.iso.1	LEV	gyrA	G668D	CRyPTIC

	CATALOGUE_VERSION	CATALOGUE_GRAMMAR	SITEID	PREDICTION
11701	v1.31	GARC1	02	S
11702	v1.31	GARC1	02	U
11703	v1.31	GARC1	02	S
11704	v1.31	GARC1	02	R
11705	v1.31	GARC1	02	R
11706	v1.31	GARC1	02	R
11707	v1.31	GARC1	02	S
11708	v1.31	GARC1	02	S
11709	v1.31	GARC1	02	S
11710	v1.31	GARC1	02	S
11711	v1.31	GARC1	02	U
11712	v1.31	GARC1	02	S

The net result of all these predictions needs to be sorted out; that is where the PREDICTIONS table comes in. This simply has one row per sample per drug per catalogue (version).

## 4.5 PREDICTIONS

```
[97]: PREDICTIONS=pandas.read_pickle(TABLES_PATH+"PREDICTIONS.pkl.gz")
PREDICTIONS[:3]
```

```
[97]:
```

UNIQUEID	CATALOGUE_GRAMMAR	SITEID \	DRUG	CATALOGUE_NAME	CATALOGUE_VERSION	
site.02.subj.0958.lab.22A197.iso.1	CFZ	CRyPTIC	v1.31	GARC1		
02						
	RIF	CRyPTIC	v1.31	GARC1		
02						
	MXF	CRyPTIC	v1.31	GARC1		
02						

UNIQUEID	CATALOGUE_GRAMMAR	PREDICTION \	DRUG	CATALOGUE_NAME	CATALOGUE_VERSION	
site.02.subj.0958.lab.22A197.iso.1	CFZ	CRyPTIC	v1.31	GARC1		
S						
	RIF	CRyPTIC	v1.31	GARC1		
R						
	MXF	CRyPTIC	v1.31	GARC1		
S						

DEFAULT\_CATALOGUE

UNIQUEID	DRUG	CATALOGUE_NAME	CATALOGUE_VERSION	CATALOGUE_GRAMMAR
site.02.subj.0958.lab.22A197.iso.1	CFZ	CRyPTIC	v1.31	GARC1
True				
	RIF	CRyPTIC	v1.31	GARC1
True				
	MXF	CRyPTIC	v1.31	GARC1
True				

If we pull out the rows for the sample sample as above

```
[98]: PREDICTIONS.reset_index(inplace=True)
PREDICTIONS.loc[(PREDICTIONS.UNIQUEID=="site.02.subj.0914.lab.22A148.iso.1") &
↳ (PREDICTIONS.DRUG.isin(['LEV', 'MXF', 'OFX']))]
```

```
[98]:
```

	UNIQUEID	DRUG	CATALOGUE_NAME	\
3956	site.02.subj.0914.lab.22A148.iso.1	LEV	CRyPTIC	
3963	site.02.subj.0914.lab.22A148.iso.1	MXF	CRyPTIC	
3966	site.02.subj.0914.lab.22A148.iso.1	OFX	CRyPTIC	

	CATALOGUE_VERSION	CATALOGUE_GRAMMAR	SITEID	PREDICTION	DEFAULT_CATALOGUE
3956	v1.31	GARC1	02	R	True
3963	v1.31	GARC1	02	R	True
3966	v1.31	GARC1	02	R	True

So they are all predicted to be resistant to the fluoroquinolones.

The logic is what you expect; if there is >0 rows in **EFFECTS** for a drug that predict resistance, then the sample is predicted to R regardless of what the other rows predict. If there are no rows in **EFFECTS** that predict resistance but >0 rows with **PREDICTION**=='U', then the sample is predicted U for that sample regardless of the other rows. If there are 0 rows, or >0 rows which are all predicted to be S, then the sample is predicted to be S.

## 4.6 GPI\_SNP\_DISTANCES arrays

These are not tables, but instead are (NxN) **numpy** arrays of SNP distances between each GPI sample and every other GPI sample. The array is therefore symmetric with a leading diagonal of zeros.

For convenience, the labels are stored in a separate **numpy** array.

If you are not familiar with **numpy** please get in touch; keeping them in this form should save you a lot of time and effort.

```
[99]: labels=numpy.load('GPI_SNP_DISTANCES_LABELS.npy')
labels
```

```
[99]: array(['site.QA.subj.N0004.lab.N0004.iso.1',
'site.QA.subj.N0031.lab.N0031.iso.1',
```

```
'site.QA.subj.N0052.lab.N0052.iso.1', ...,
'site.21.subj.004.lab.MR253325K.iso.1',
'site.21.subj.005.lab.MR311341D.iso.1',
'site.21.subj.006.lab.MR304907G.iso.1'], dtype='<U57')
```

Note that the samples with SITEID=='QA' are the 17 Comas samples. 15211+17=15228

```
[100]: len(labels)
```

```
[100]: 15228
```

```
[101]: distances=numpy.load('GPI_SNP_DISTANCES_VALUES.npy')
distances
```

```
[101]: array([[ 0, 1118, 1140, ..., 1198, 1157, 407],
 [1118, 0, 829, ..., 1138, 847, 1108],
 [1140, 829, 0, ..., 1166, 583, 1125],
 ...,
 [1198, 1138, 1166, ..., 0, 1196, 1192],
 [1157, 847, 583, ..., 1196, 0, 1157],
 [407, 1108, 1125, ..., 1192, 1157, 0]], dtype=int16)
```

Selecting the distances from one sample is easy with numpy fancy indexing. For example, to find all the SNP distances to the first sample we simply first create a array of Booleans telling us which column has the name of the first sample (the first one, funnily enough).

```
[102]: labels=='site.QA.subj.N0004.lab.N0004.iso.1'
```

```
[102]: array([ True, False, False, ..., False, False, False])
```

Then we index the distances using that Boolean array

```
[103]: d=distances[labels=='site.QA.subj.N0004.lab.N0004.iso.1']
d
```


```
[103]: array([[ 0, 1118, 1140, ..., 1198, 1157, 407]], dtype=int16)
```

and we can simply calculate some statistics and plot a histogram

```
[104]: numpy.average(d)
```

```
[104]: 1122.3733254531128
```

```
[105]: fig,axis=plt.subplots(1,1,figsize=(8,5))
a=axis.hist(d.flatten(),bins=numpy.arange(0,2500,100))
```



Let's fish out all theGPI samples that are within 500 SNPs of this QA strain. This is where the power and simplicity of the `numpy` fancy indexing comes into its own!

The condition produces an array of Booleans..

```
[106]: close_samples=d<400
       close_samples
```

```
[106]: array([[ True, False, False, ..., False, False, False]])
```

..which we can then use to pull out the distances

```
[107]: d[close_samples]
```

```
[107]: array([ 0, 380, 346, ..., 344, 349, 335], dtype=int16)
```

..and which samples they came from

```
[108]: labels[close_samples[0]]
```

```
[108]: array(['site.QA.subj.N0004.lab.N0004.iso.1',
             'site.QA.subj.N0054.lab.N0054.iso.1',
             'site.02.subj.0370.lab.222044-14.iso.1', ...,
             'site.20.subj.SCH8604399.lab.YA00134971.iso.1',
             'site.21.subj.002.lab.MM011229E.iso.1',
             'site.21.subj.003.lab.MR428686L.iso.1'], dtype='<U57')
```

Finally, we can use this array to pull the rows in the GENOMES tables for these samples and then look which lineages they belong to

```
[109]: GENOMES.loc[GENOMES.index.isin(labels[close_samples[0]])].LINEAGE_NAME.  
       ↪value_counts()
```

```
[109]: Lineage 3      1508  
       Lineage 4       37  
       Lineage 2        1  
       Lineage 6        0  
       Lineage 5        0  
       Lineage 1        0  
              0  
       Name: LINEAGE_NAME, dtype: int64
```

Let's try something more ambitious and calculate the SNP distributions for the distances between each and every set of lineages. (This will take about 30 sec to run)

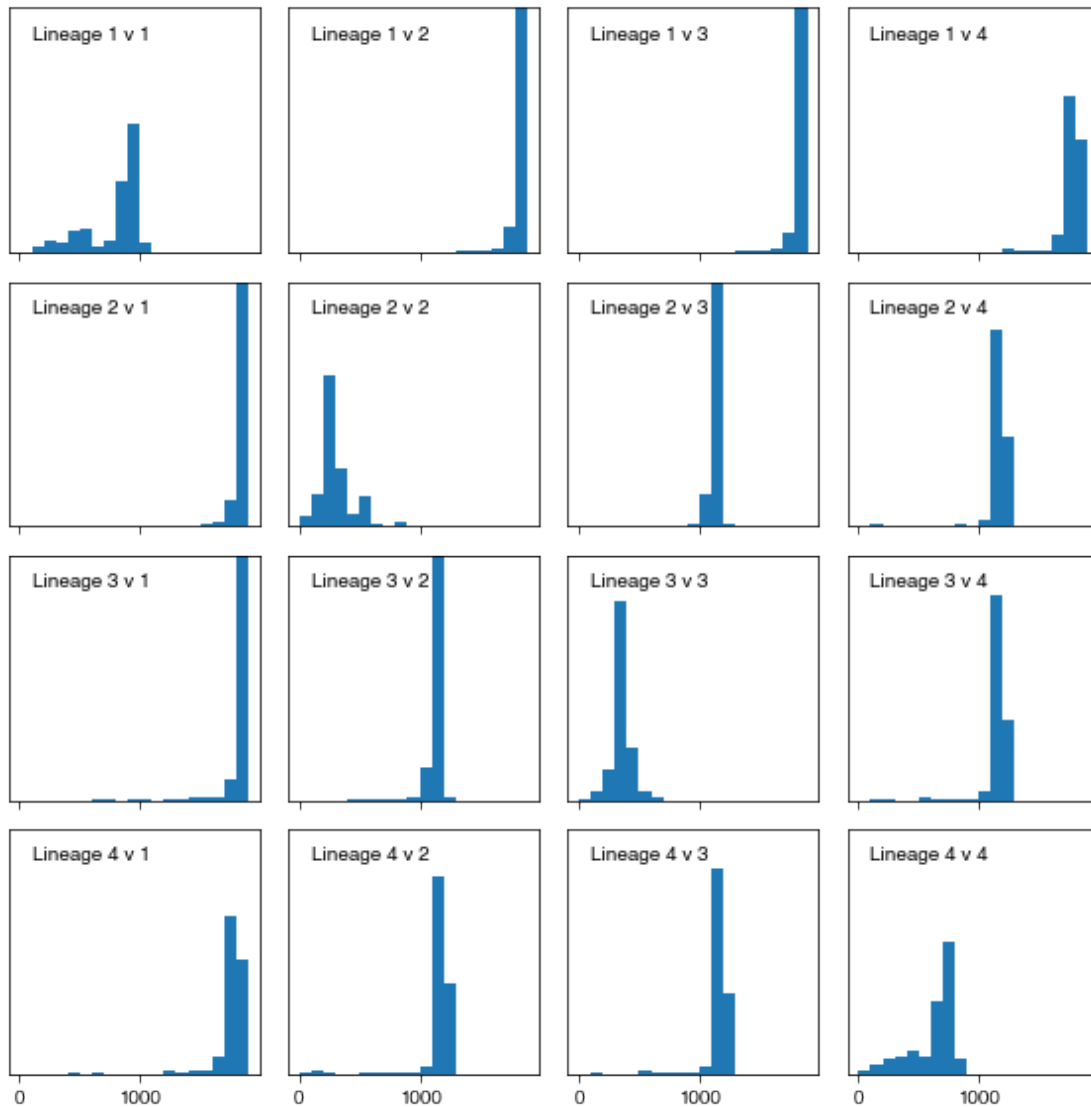
```
[110]: l={}
       for lineage in [1,2,3,4]:
           l[lineage]=numpy.isin(labels,list(GENOMES.loc[GENOMES.
           ↪LINEAGE_NAME=='Lineage '+str(lineage)].index))

       d={}
       for lineage1 in [1,2,3,4]:
           for lineage2 in [1,2,3,4]:
               d[(lineage1,lineage2)]=distances[numpy.ix_(l[lineage1],l[lineage2])]

       fig,axes=plt.subplots(4,4,sharex=True,tight_layout=True,figsize=(8,8))

       for lineage1 in [1,2,3,4]:
           for lineage2 in [1,2,3,4]:
               axes[lineage1-1,lineage2-1].set_ylim([0,0.008])
               axes[lineage1-1,lineage2-1].axes.get_yaxis().set_visible(False)
               axes[lineage1-1,lineage2-1].hist(d[(lineage1,lineage2)].
               ↪flatten(),bins=numpy.arange(0,2000,100),density=True)
               axes[lineage1-1,lineage2-1].text(100,0.007,"Lineage "+str(lineage1)+" v_⊥
               ↪"+str(lineage2),horizontalalignment='left')
```





As expected, all samples that belong to a lineage are more similar to one another (generally  $< 1000$  SNPs) than when samples belong to two different lineages are compared. The exception is when comparing Lineages 2 and 3 which are more alike than any of the other lineage pairs.

This result provides some comfort that (a) the SNP distance calculated by the Clockwork resequencing process is correct and (b) the lineages called by SNP-IT are also good.

## 5 Analysis Examples

### 5.1 rpoB\_S450L MIC distribution

```
[111]: PHENOTYPES=pandas.read_pickle(TABLES_PATH+"UKMYC_PHENOTYPES.pkl.gz")

PHENOTYPES.reset_index(inplace=True)

PHENOTYPES=PHENOTYPES.loc[(PHENOTYPES.DRUG=="RIF") &\
                           (PHENOTYPES.PLATEDESIGN=="UKMYC5") &\
                           (PHENOTYPES.PHENOTYPE_QUALITY=="HIGH") &\
                           (PHENOTYPES.DILUTION>0)]

PHENOTYPES.set_index(["UNIQUEID"],inplace=True,verify_integrity=True)

PHENOTYPES[:3]
```

```
[111]:
```

	DRUG	PLATEDESIGN	BELONGS_GPI	\
UNIQUEID				
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	RIF	UKMYC5	True	
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	RIF	UKMYC5	True	
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	RIF	UKMYC5	True	

	SITEID	DILUTION	\
UNIQUEID			
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	05	8.0	
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	06	8.0	
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	05	2.0	

	PHENOTYPE_QUALITY	READINGDAY	\
UNIQUEID			
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	HIGH	14	
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	HIGH	14	
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	HIGH	14	

	PRIMARY_DILUTION	PRIMARY_METHOD	\
UNIQUEID			
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	8.0	VZ	
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	8.0	VZ	
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	2.0	VZ	

	AMYGDA_DILUTION	\
UNIQUEID		
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	1.0	
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	8.0	
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	2.0	

UNIQUEID	BASHTHEBUG_DILUTION \
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	8.0
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	8.0
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	2.0

UNIQUEID	BASHTHEBUGPRO_DILUTION \
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	NaN
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	NaN
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	NaN

UNIQUEID	PHENOTYPE_DESCRIPTION \
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	VZ,BB AGREE
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	VZ,IM AGREE
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	VZ,IM AGREE

UNIQUEID	BASHTHEBUG_NUMBER_CLASSIFICATIONS \
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	16.0
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	11.0
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	11.0

UNIQUEID	MIC	LOG2MIC	BINARY_PHENOTYPE
site.05.subj.PSLM-0812.lab.SLM-072.iso.1	>4	3.00	R
site.06.subj.SSM_0145-14.lab.06MIL0268.iso.1	>4	3.00	R
site.05.subj.PTAN-0394.lab.TAN-650.iso.1	0.12	-3.06	S

```
[114]: MUTATIONS=pandas.read_pickle(TABLES_PATH+"MUTATIONS_SAMPLE.pkl.gz")

MUTATIONS.reset_index(inplace=True)

MUTATIONS=MUTATIONS.loc[(MUTATIONS.MUTATION=="S450L") & (MUTATIONS.
↪GENE=="rpoB")]

MUTATIONS.set_index(["UNIQUEID"],inplace=True,verify_integrity=True)

MUTATIONS[:3]
```

UNIQUEID	GENE	MUTATION	POSITION \
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	rpoB	S450L	450.0
site.05.subj.LI2076709.lab.15277_3_50.iso.1	rpoB	S450L	450.0

site.05.subj.PSLM-0779.lab.SLM-034.iso.1	rpoB	S450L	450.0	
--	------	-------	-------	--

	AMINO_ACID_NUMBER	GENOME_INDEX	\	
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	450.0	NaN		
site.05.subj.LI2076709.lab.15277_3_50.iso.1	450.0	NaN		
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	450.0	NaN		

	NUCLEOTIDE_NUMBER	REF	ALT	\
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	NaN	tcg	ttg	
site.05.subj.LI2076709.lab.15277_3_50.iso.1	NaN	tcg	ttg	
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	NaN	tcg	ttg	

	IS_SNP	IS_INDEL	IN_CDS	\
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	True	False	True	
site.05.subj.LI2076709.lab.15277_3_50.iso.1	True	False	True	
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	True	False	True	

	IN_PROMOTER	IS_SYNONYMOUS	\	
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	False	False		
site.05.subj.LI2076709.lab.15277_3_50.iso.1	False	False		
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	False	False		

	IS_NONSYNONYMOUS	IS_HET	\	
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	True	False		
site.05.subj.LI2076709.lab.15277_3_50.iso.1	True	False		
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	True	False		

	IS_NULL	IS_FILTER_PASS	\	
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	False	True		
site.05.subj.LI2076709.lab.15277_3_50.iso.1	False	True		
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	False	True		

	ELEMENT_TYPE	MUTATION_TYPE	\	
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	GENE	AAM		
site.05.subj.LI2076709.lab.15277_3_50.iso.1	GENE	AAM		
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	GENE	AAM		


  

	INDEL_LENGTH	INDEL_1	INDEL_2	\
UNIQUEID				
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	NaN			

```
site.05.subj.LI2076709.lab.15277_3_50.iso.1      NaN
site.05.subj.PSLM-0779.lab.SLM-034.iso.1         NaN
```

UNIQUEID	SITEID	NUMBER_NUCLEOTIDE_CHANGES
site.05.subj.LR-2264.lab.FN-00887-17.iso.1	05	1
site.05.subj.LI2076709.lab.15277_3_50.iso.1	05	1
site.05.subj.PSLM-0779.lab.SLM-034.iso.1	05	1

```
[115]: MUTATIONS=MUTATIONS[["ALT"]]
df=PHENOTYPES.join(MUTATIONS,how="inner")
a=df.DILUTION.hist()
```



So as expected, the majority of samples with an `rpoB@S450L` mutation have growth in all wells on the UKMYC5 plate.